

# ITC 5104 Database Design and SQL

Lecture 1

Chapter 2 Oracle 12c: SQL

Basic SQL SELECT Statements

# Objectives

- Create initial database
- Identify keywords, mandatory clauses, and optional clauses in a SELECT statement
- Select and view all, one or multiple columns of a table
- Use a column alias to clarify contents of a column
- Perform basic arithmetic operations in a SELECT clause
- Remove duplicate items using either DISTINCT or UNIQUE keywords
- Some basic settings for SQL Developer

# Create Initial Database

- For this course we will use the SQL Developer interface as a client software tool
- SQL Developer will be used to connect to the Oracle 12c database residing on Dilbert
- For those of you with the textbook you can refer to Appendix B for an overview of SQL Developer (is based on previous version)
- These steps were outlined in both Lab Exercise 1 and the notes on installing SQL Developer on your own system

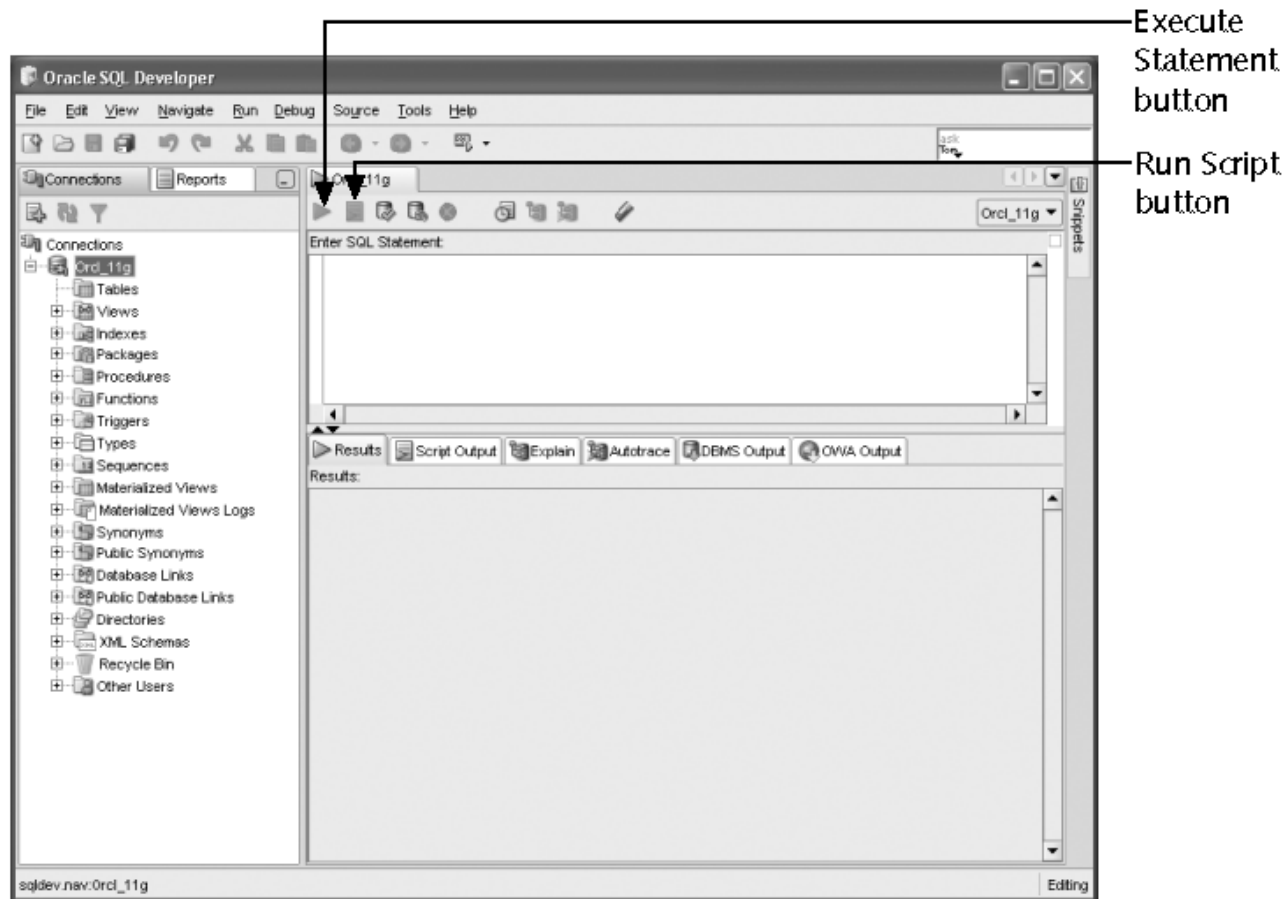
# Some Facts About Data

- 2.5 quintillion bytes of data generated each day 2.5 exabytes
- This equates to about 1 Zettabytes per year or (1000 Exabyte)
- Facebook scans roughly **105 terabytes** of data each half hour

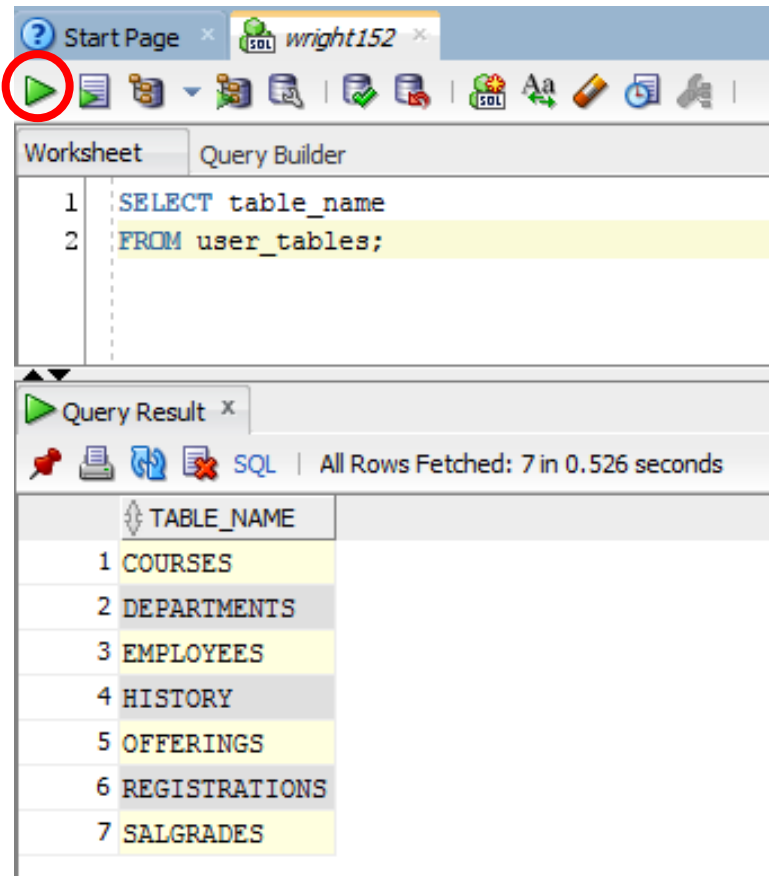
# Some Facts About Data

- Walmart handles 2.5 petabytes every hour in transactions
  - Walmart building cloud database to process 2.5 petabytes per hour
- Business data doubles every 1.2 years
  - As of July 2017 over 1.2 Billion active users for Gmail

# SQL Developer Interface

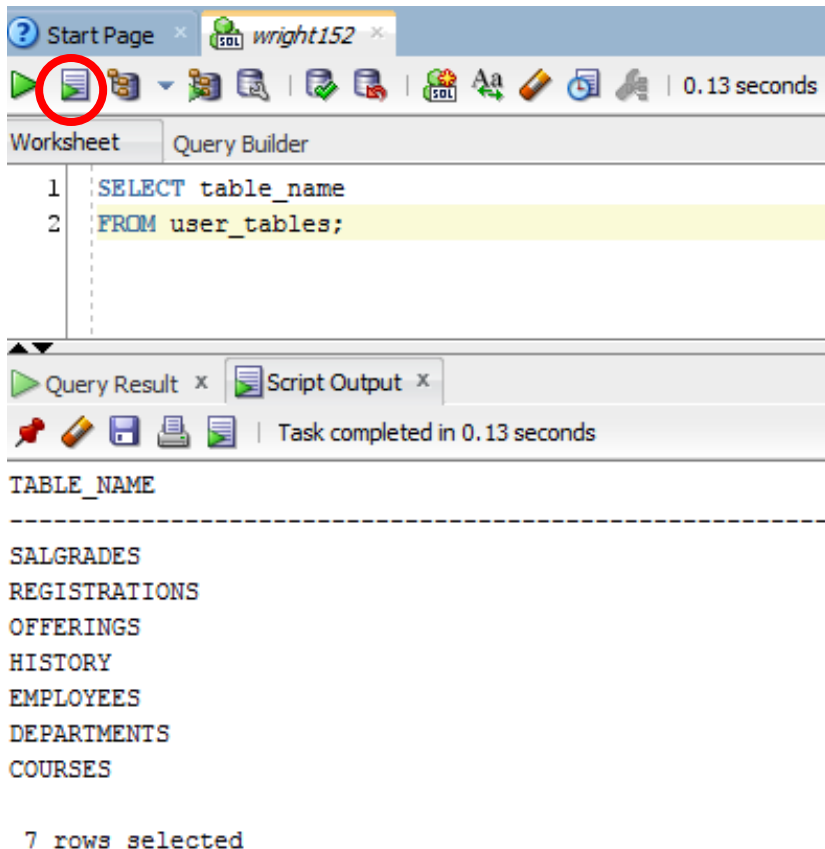


# What Tables Do I Own?



- user\_tables is a data dictionary view that will display the tables you own
- To execute this SQL statement the Execute Statement button is clicked
- You can also press the Ctrl+Enter keys to execute a statement
- You have to have your cursor in the statement you wish to execute, other wise nothing happens

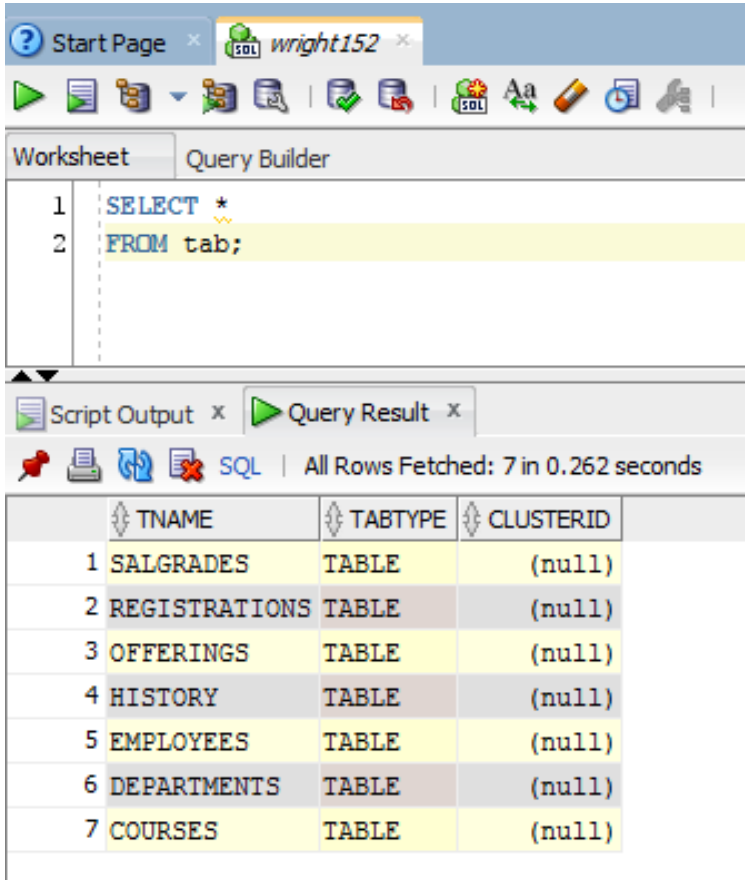
# What Tables Do I Own?



- The output is the same, just in a different format
- This output was produced by clicking the Run Script button or by pressing the F5 function key
- This produces a text based output not as structured as the previous method
- The window in this case does not clear each time the button is pressed, instead it will accumulate and show all outputs that have been generated



# What Tables Do I Own?



The screenshot shows a SQL query editor with the following query:

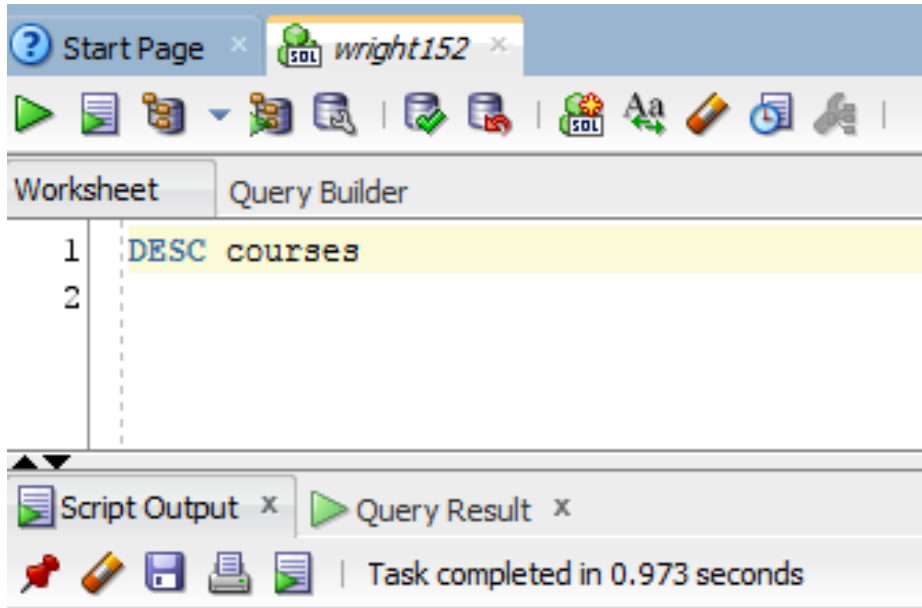
```
1 SELECT *
2 FROM tab;
```

Below the query editor, the 'Query Result' pane displays a table with 7 rows and 3 columns: TNAME, TABTYPE, and CLUSTERID. The rows list tables: SALGRADES, REGISTRATIONS, OFFERINGS, HISTORY, EMPLOYEES, DEPARTMENTS, and COURSES, all of type TABLE.

	TNAME	TABTYPE	CLUSTERID
1	SALGRADES	TABLE	(null)
2	REGISTRATIONS	TABLE	(null)
3	OFFERINGS	TABLE	(null)
4	HISTORY	TABLE	(null)
5	EMPLOYEES	TABLE	(null)
6	DEPARTMENTS	TABLE	(null)
7	COURSES	TABLE	(null)

- This is an alternate table called TAB
- TAB is a pseudo table, you will not see it but you can use it
- Its purpose is to display the table objects and other objects called VIEWS
- Will also use to see any dropped tables you may have in your schema

# DESCRIBE the Structure of a Table

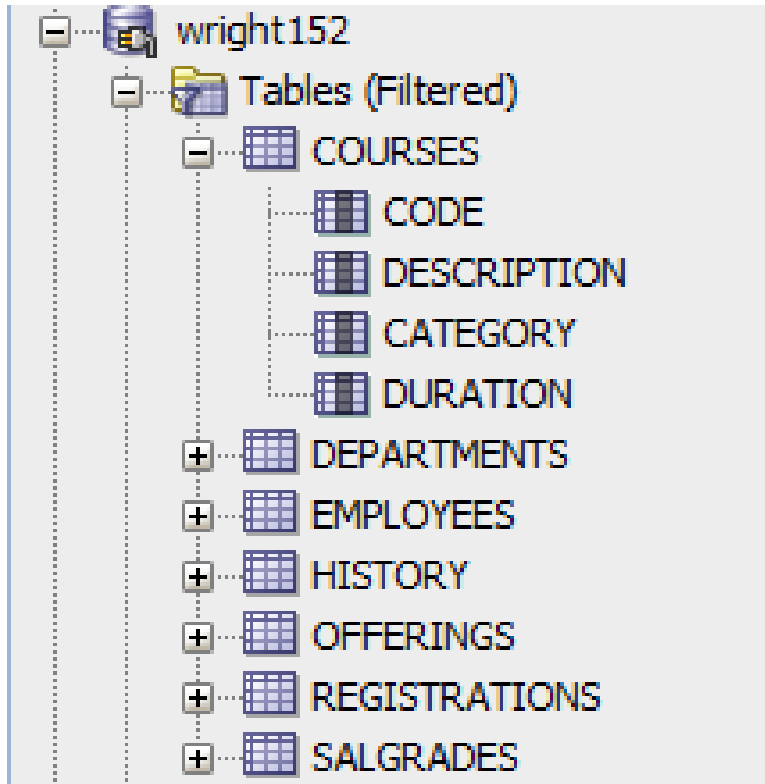


The screenshot shows a database management tool interface. The top pane, labeled 'Query Builder', contains the command `DESC courses` in a yellow-highlighted cell. The bottom pane, labeled 'Query Result', displays the output of the command. It shows the table structure for 'courses' with columns: Name, Null, and Type. The data is as follows:

Name	Null	Type
CODE	NOT NULL	VARCHAR2 (6)
DESCRIPTION	NOT NULL	VARCHAR2 (30)
CATEGORY	NOT NULL	CHAR (3)
DURATION	NOT NULL	NUMBER (2)

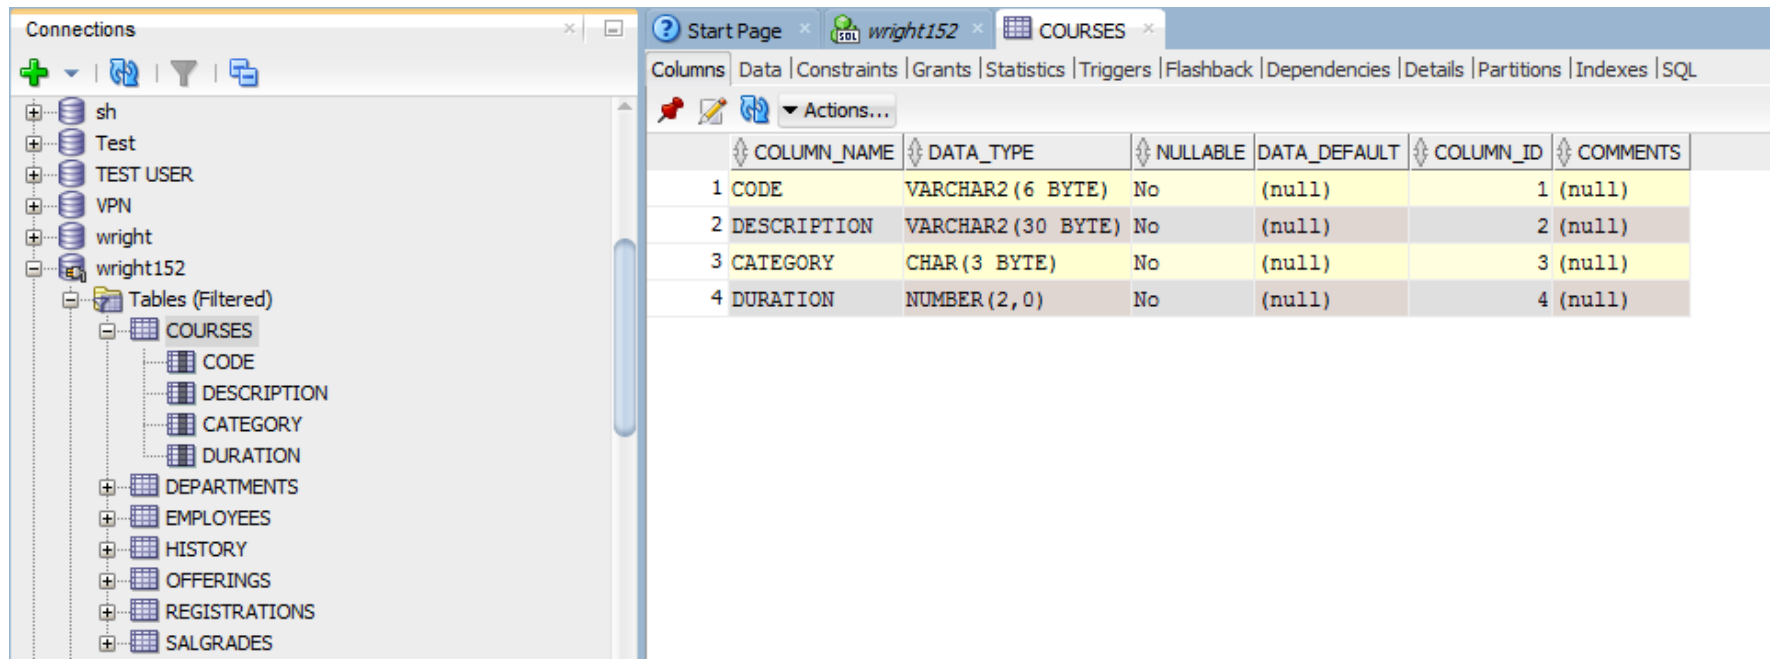
- The DESC or DESCRIBE command followed by a table name will show the structure of that table
- The structure is the names of the columns and their data types with sizes
- The NULL column shows any columns that have been defined as NOT NULL

# Viewing Tables in SQL Developer



- In SQL you can expand the tree where it says Tables, this is on the left side of the window
- The tables you have in your database account will be displayed
- You can click on the actual table name to see details of the table

# Viewing Tables in SQL Developer



- When you click on a table name the columns in the table are shown in the tree and details of the table appear to the right in the upper pane
- You will also notice a series of Tabs above the output, you can click these to see various aspects of the table

# SELECT Statement

- The majority of the SQL operations performed on a database are done with the **SELECT** statement
- The **SELECT** statement allows a user to retrieve data from existing tables
- Note that **SELECT does not change data**, it simply retrieves it from the database
- The user can ask for all of the fields and records in a table or can request only certain fields and records, up to you as the user
- The **SELECT** statement asks the database a question, also known as a **query**

# SELECT Statement

- After querying the database, the results are displayed
- What is displayed is the answer to the question (query) asked by the user
- Keywords begin a section of a **SELECT** statement are called a **clause**:
  - **SELECT clause**
  - **FROM clause**
  - **WHERE clause**
  - **ORDER BY clause**
  - **There are more of them**

# SELECT Statement Syntax

```
SELECT  [DISTINCT | UNIQUE] (*, columnname [ AS alias], ...)  
      FROM      tablename  
      [WHERE     condition]  
      [GROUP BY  group_by_expression]  
      [HAVING    group_condition]  
      [ORDER BY  columnname];
```

This is a syntax diagram for the SELECT statement. We will discuss all of these pieces and even a few that are not here over the course of the semester

We will make reference to this syntax diagram again and again as we build the SELECT statement

# SELECT Statement Syntax Style

- The other capitalized components in the **SELECT** statement also begin clauses
- The only required clauses are **SELECT** and **FROM**
- The asterisk character is used to indicate that all columns available are to be displayed
- To select the entire table, the syntax is:

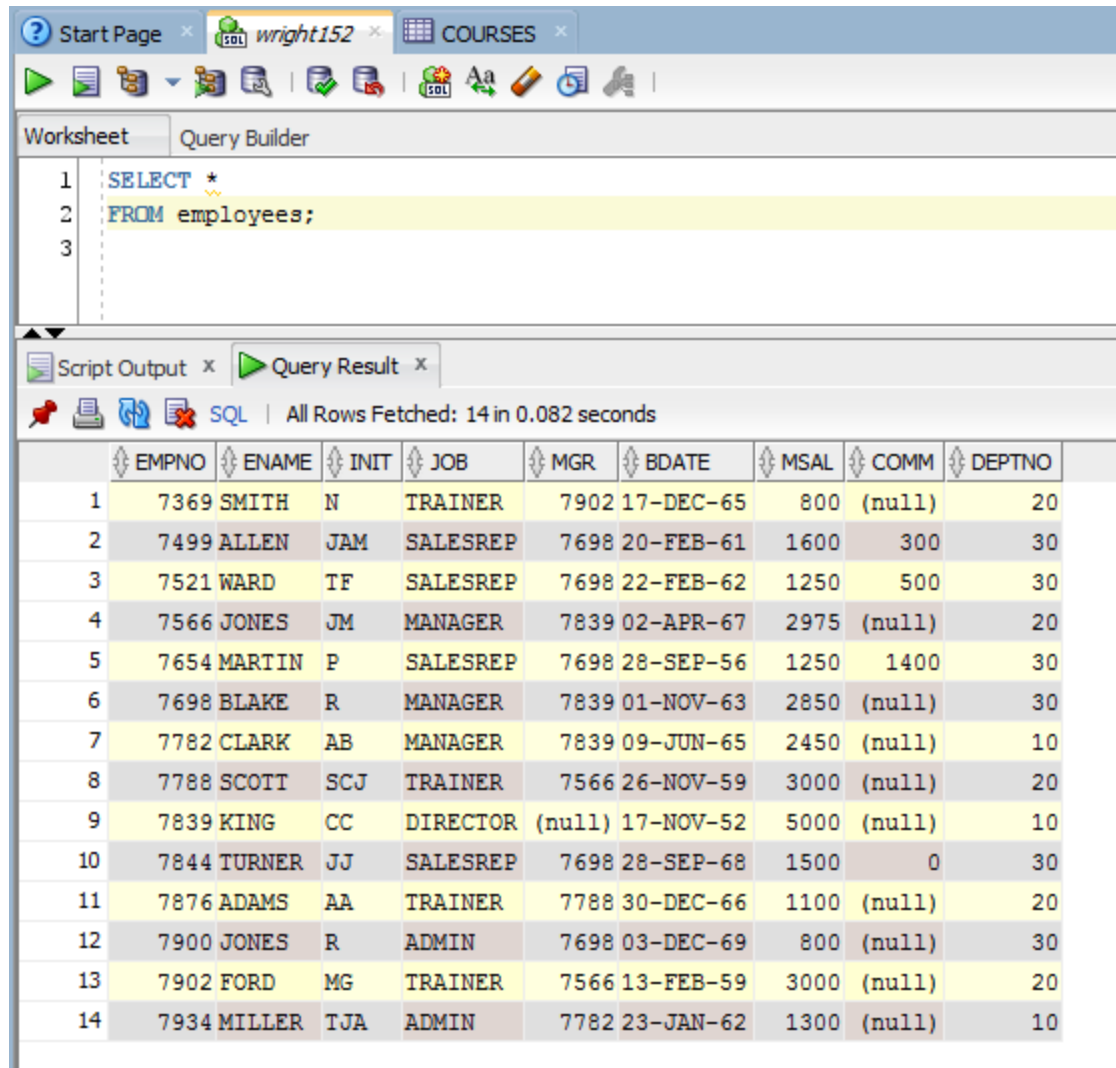
**SELECT \***

**FROM *table\_name*;**

- By default, column headings displayed in a query are capitalized
- \*\*\* It is best to enter your SQL command over several lines, beginning each line with a keyword \*\*\*
- \*\*\* Please enter query over multiple lines, it is not a contest to see who can use the fewest lines, make it clear and readable, and easier to debug \*\*\*



# Selecting all Rows in a Table



The screenshot shows a SQL query editor interface. The top bar includes tabs for 'Start Page', 'wright152', and 'COURSES'. Below the tabs is a toolbar with various icons. The main area is divided into 'Worksheet' and 'Query Builder' tabs. The 'Worksheet' tab is active, showing a SQL query:

```
1 SELECT *
2 FROM employees;
3
```

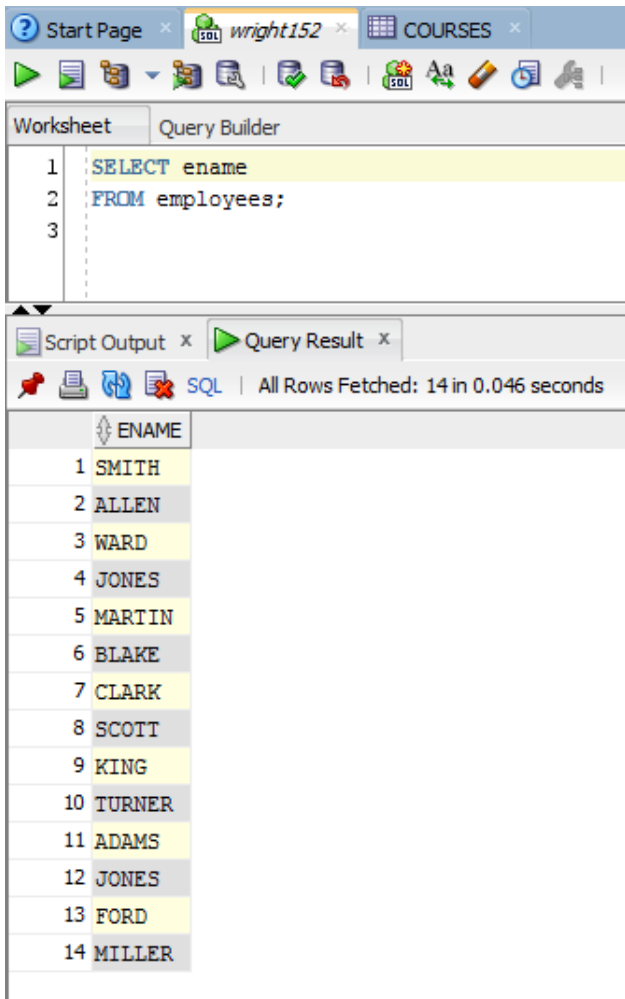
Below the query editor, there is a 'Script Output' and 'Query Result' tab. The 'Query Result' tab is active, displaying the results of the query. The status bar indicates 'All Rows Fetched: 14 in 0.082 seconds'.

	EMPNO	ENAME	INIT	JOB	MGR	BDATE	MSAL	COMM	DEPTNO
1	7369	SMITH	N	TRAINER	7902	17-DEC-65	800	(null)	20
2	7499	ALLEN	JAM	SALESREP	7698	20-FEB-61	1600	300	30
3	7521	WARD	TF	SALESREP	7698	22-FEB-62	1250	500	30
4	7566	JONES	JM	MANAGER	7839	02-APR-67	2975	(null)	20
5	7654	MARTIN	P	SALESREP	7698	28-SEP-56	1250	1400	30
6	7698	BLAKE	R	MANAGER	7839	01-NOV-63	2850	(null)	30
7	7782	CLARK	AB	MANAGER	7839	09-JUN-65	2450	(null)	10
8	7788	SCOTT	SCJ	TRAINER	7566	26-NOV-59	3000	(null)	20
9	7839	KING	CC	DIRECTOR	(null)	17-NOV-52	5000	(null)	10
10	7844	TURNER	JJ	SALESREP	7698	28-SEP-68	1500	0	30
11	7876	ADAMS	AA	TRAINER	7788	30-DEC-66	1100	(null)	20
12	7900	JONES	R	ADMIN	7698	03-DEC-69	800	(null)	30
13	7902	FORD	MG	TRAINER	7566	13-FEB-59	3000	(null)	20
14	7934	MILLER	TJA	ADMIN	7782	23-JAN-62	1300	(null)	10

# Selecting a Single Column

- In most cases, you will only want to see the data in certain columns
- This may be because some data is sensitive and you do not want others to see it, or the results wrap because of the number of columns being displayed
- Could also be that there is just too much data to see and you do not need to see all the data
- Selecting specific columns in a **SELECT** statement is called *projection*. You may select one or many columns, as desired

# Selecting a Single Column



The screenshot shows a SQL query editor with a 'Query Builder' tab. The query is as follows:

```
1 SELECT ename
2 FROM employees;
3
```

Below the query editor, the 'Query Result' tab is active, displaying the results of the query. The results are shown in a table with 14 rows, each containing an employee's name. The table has a single column labeled 'ENAME'.

	ENAME
1	SMITH
2	ALLEN
3	WARD
4	JONES
5	MARTIN
6	BLAKE
7	CLARK
8	SCOTT
9	KING
10	TURNER
11	ADAMS
12	JONES
13	FORD
14	MILLER

- Only the ename column is requested in the query result
- This query shows all of the ename values in the EMPLOYEES table

# Writing SQL Statements – Syntax Rules

- **SQL** statements are **not** case sensitive
- They can be entered on one or many lines
- Keywords cannot be split across lines or abbreviated
- Clauses are usually placed on separate lines for readability and ease of editing
- Tabs and indents make code more readable
- The preferred style is for Keywords to be entered in uppercase, while all other words such as table names and columns to be entered in lowercase (companies usually have coding standards)
- SQL statements end with a semi-colon (;)
- Use of the ; is shown on the previous slide
- **If you are executing only a **single command** in SQL Developer you can omit the semi-colon, good practice to include the semi-colon though**

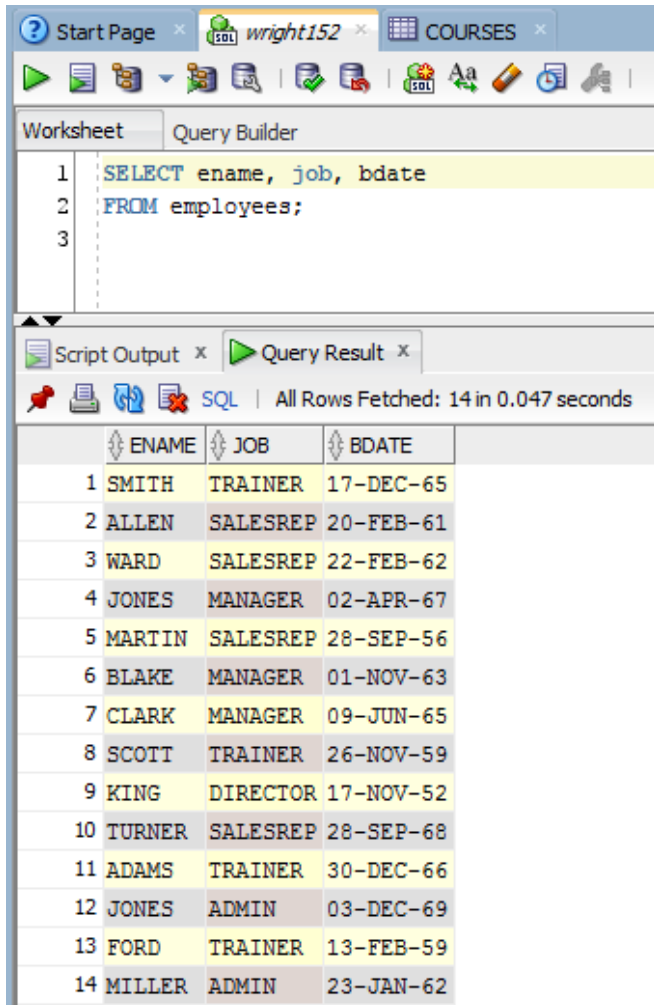
# Writing SQL Statements – Best Practices

- Based on the previous information, will the following SQL statements produce the same results as the previous slide?
  1. `SELECT ENAME FROM EMPLOYEES;`
  2. `select ename from employees;`
  3. `SELECT ename FROM employees;`
  4. `SELECT ename  
FROM employees;`
- Which statement do I expect to use in this case?

# Selecting Multiple Columns From a Table

- The **SELECT** statement can also be used to retrieve multiple columns from a table
- A comma is used to separate the column names, a space is not necessary but is used to improve readability
- The order the columns are listed is the order in which the columns will display

# Selecting Multiple Columns From a Table



The screenshot shows a database query tool interface. At the top, there are tabs for 'Start Page', 'wright152', and 'COURSES'. Below the tabs is a toolbar with various icons. The main area is divided into two sections: 'Worksheet' and 'Query Builder'. The 'Worksheet' section contains a SQL query:

```
1 SELECT ename, job, bdate
2 FROM employees;
3
```

The 'Query Result' section shows the results of the query. It includes a status bar that says 'All Rows Fetched: 14 in 0.047 seconds'. Below the status bar is a table with three columns: ENAME, JOB, and BDATE. The table contains 14 rows of data.

	ENAME	JOB	BDATE
1	SMITH	TRAINER	17-DEC-65
2	ALLEN	SALESREP	20-FEB-61
3	WARD	SALESREP	22-FEB-62
4	JONES	MANAGER	02-APR-67
5	MARTIN	SALESREP	28-SEP-56
6	BLAKE	MANAGER	01-NOV-63
7	CLARK	MANAGER	09-JUN-65
8	SCOTT	TRAINER	26-NOV-59
9	KING	DIRECTOR	17-NOV-52
10	TURNER	SALESREP	28-SEP-68
11	ADAMS	TRAINER	30-DEC-66
12	JONES	ADMIN	03-DEC-69
13	FORD	TRAINER	13-FEB-59
14	MILLER	ADMIN	23-JAN-62

- The **ename, job and bdate** columns appear in the result table
- The ename was specified first in the SELECT CLAUSE so it appears first followed by the job, then by the bdate column
- A comma separates the two column names
- The statement is terminated with a semicolon
- The columns display in the order they are specified

# Operations Within the SQL Statement

- We will now look at some other operations that can be done with the SELECT statement:
  - Column aliases
  - Arithmetic operations
  - Elimination of duplicate output
  - Display rows on multiple lines



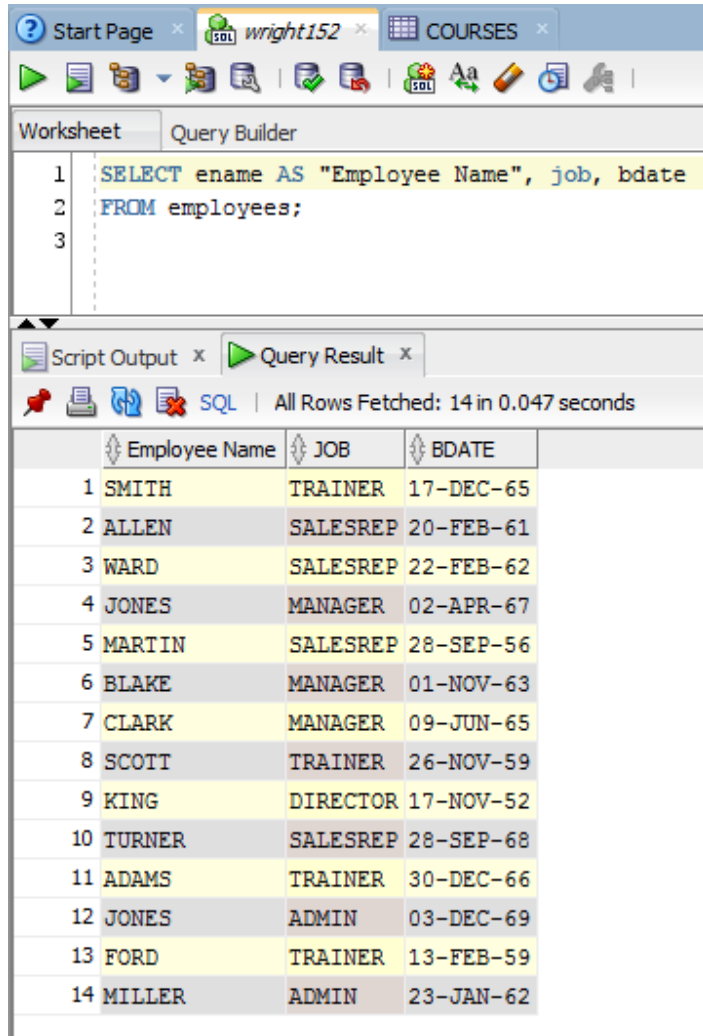
# Using Column Aliases

- In some cases, the column (attribute) name may not be a good descriptor of the data displayed
- To better describe the data listed, you can substitute a **column alias** for the column name in the results of a query
- To use a column alias, place the name of the alias beside the column name in the query

# Using Column Aliases

- The optional keyword **AS** can be included to indicate that what follows is the column alias. It distinguishes the column name from the column alias
- If the alias name contains any spaces or special symbols, or if you want to maintain the case of the alias, you must enclose the alias in *double quotation marks*

# Using Column Aliases



The screenshot shows a SQL query editor with a query window and a query result window. The query window contains the following SQL code:

```
1 SELECT ename AS "Employee Name", job, bdate
2 FROM employees;
3
```

The query result window displays the results of the query, showing 14 rows of data. The columns are labeled Employee Name, JOB, and BDATE. The data is as follows:

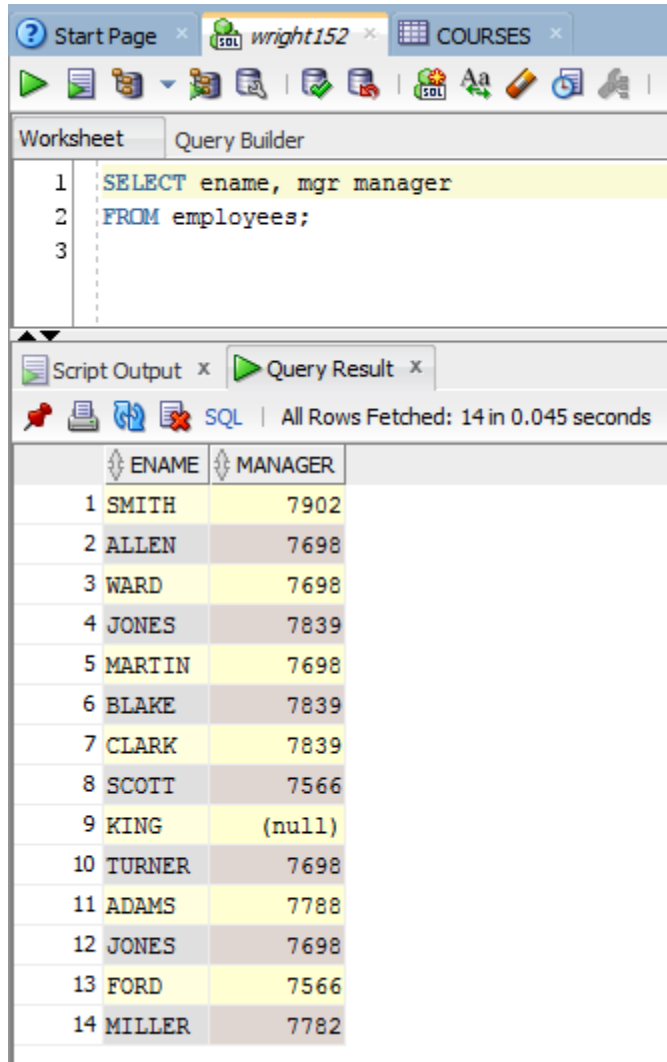
	Employee Name	JOB	BDATE
1	SMITH	TRAINER	17-DEC-65
2	ALLEN	SALESREP	20-FEB-61
3	WARD	SALESREP	22-FEB-62
4	JONES	MANAGER	02-APR-67
5	MARTIN	SALESREP	28-SEP-56
6	BLAKE	MANAGER	01-NOV-63
7	CLARK	MANAGER	09-JUN-65
8	SCOTT	TRAINER	26-NOV-59
9	KING	DIRECTOR	17-NOV-52
10	TURNER	SALESREP	28-SEP-68
11	ADAMS	TRAINER	30-DEC-66
12	JONES	ADMIN	03-DEC-69
13	FORD	TRAINER	13-FEB-59
14	MILLER	ADMIN	23-JAN-62

- The column heading ename has been replaced with the alias Employee Name
- Notice JOB and BDATE, are displayed in upper case still
- Since “Employee Name” is enclosed in double-quotations it displays in the case specified as well I am allowed to use the space character
- If I wanted all column headings to appear the same how would I accomplish this?

# Using Column Aliases

- As you can see from the next slide, it is sometimes difficult to tell the column name from the alias without the **AS** keyword – retail is the column name, price is the alias!

# Using Column Aliases



The screenshot shows the SQL Developer interface. The top toolbar includes icons for Start Page, Query Builder, and various SQL operations. The main window is divided into two panes: 'Worksheet' and 'Query Builder'. The 'Worksheet' pane contains the following SQL query:

```
1 SELECT ename, mgr manager
2 FROM employees;
3
```

The 'Query Result' pane shows the results of the query. It displays a table with two columns: 'ENAME' and 'MANAGER'. The table contains 14 rows of data, including employee names and their manager IDs. The status bar indicates 'All Rows Fetched: 14 in 0.045 seconds'.

	ENAME	MANAGER
1	SMITH	7902
2	ALLEN	7698
3	WARD	7698
4	JONES	7839
5	MARTIN	7698
6	BLAKE	7839
7	CLARK	7839
8	SCOTT	7566
9	KING	(null)
10	TURNER	7698
11	ADAMS	7788
12	JONES	7698
13	FORD	7566
14	MILLER	7782

- The AS keyword sometimes clarifies that an alias was actually requested
- Here it could be confusing since MANAGER could have been an actual column and a comma was omitted in the SELECT statement
- MANAGER is the alias for MGR

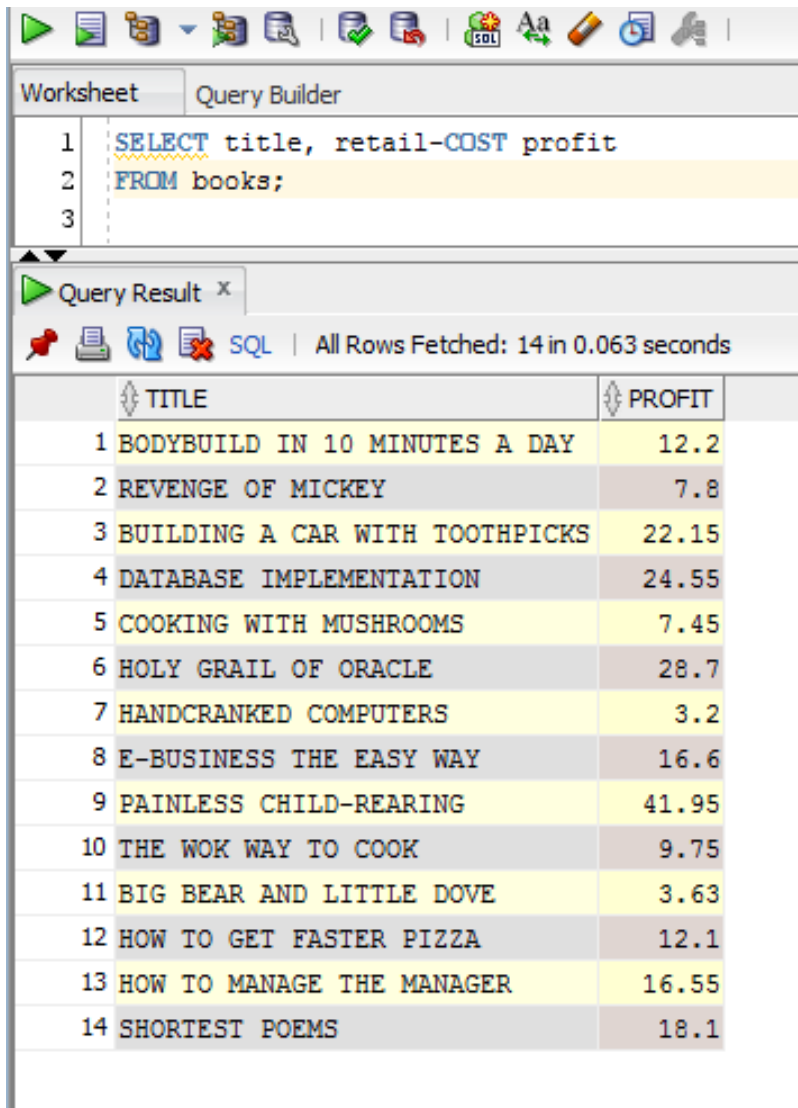
# Column Alignment

- Data for text or character fields is left-aligned
- Data for date fields is left-aligned
- Data for numeric data is right-aligned
- By default, Oracle does not display insignificant zeroes.
- If there was a value of 54.50, the zero on the right side is insignificant so it is suppressed in the result, you would only see 54.5 display
- 25 is would actually be displayed for 25.00, the zeroes are insignificant so they are suppressed in the result

# Arithmetic Expressions

- You can create expressions on **NUMBER** and **DATE** data by using arithmetic operators:
  - + Add
  - - Subtract
  - \* Multiply
  - / Divide
- **Arithmetic operators** can be used in *any clause* of a SQL statement **except** the **FROM** clause
- **Exponents** are not supported, use multiplication
- Arithmetic operations follow the standard order of operations (BEDMAS, or PEMDAS if you learned either of these acronyms)

# Using Arithmetic Operations



The screenshot shows a database query builder interface. The top toolbar contains various icons for file operations, execution, and formatting. Below the toolbar, there are two tabs: 'Worksheet' and 'Query Builder'. The 'Query Builder' tab is active, displaying a SQL query in a text area:

```
1 SELECT title, retail-COST profit
2 FROM books;
3
```

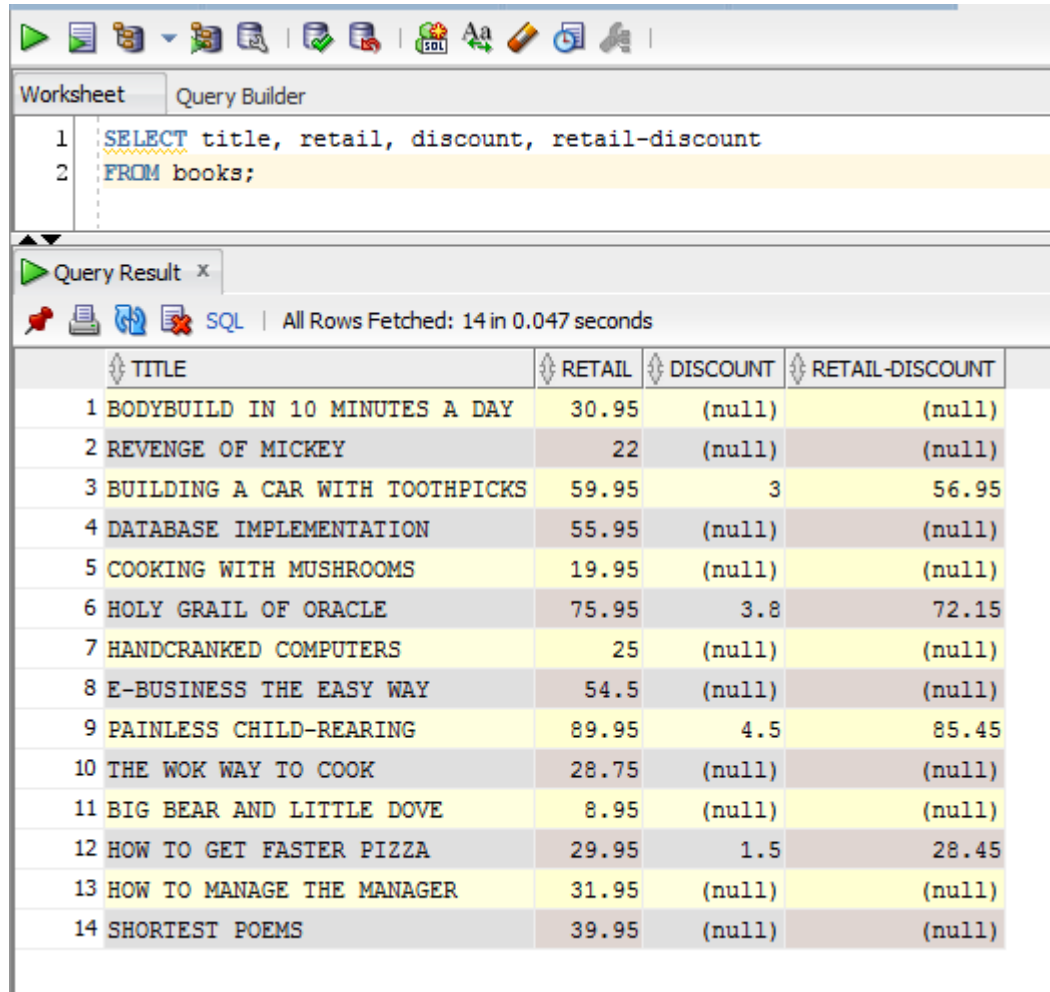
Below the query editor, there is a 'Query Result' window. It shows the execution status: 'All Rows Fetched: 14 in 0.063 seconds'. The results are displayed in a table with two columns: 'TITLE' and 'PROFIT'.

	TITLE	PROFIT
1	BODYBUILD IN 10 MINUTES A DAY	12.2
2	REVENGE OF MICKEY	7.8
3	BUILDING A CAR WITH TOOTHPICKS	22.15
4	DATABASE IMPLEMENTATION	24.55
5	COOKING WITH MUSHROOMS	7.45
6	HOLY GRAIL OF ORACLE	28.7
7	HANDCRANKED COMPUTERS	3.2
8	E-BUSINESS THE EASY WAY	16.6
9	PAINLESS CHILD-REARING	41.95
10	THE WOK WAY TO COOK	9.75
11	BIG BEAR AND LITTLE DOVE	3.63
12	HOW TO GET FASTER PIZZA	12.1
13	HOW TO MANAGE THE MANAGER	16.55
14	SHORTEST POEMS	18.1

- Notice the use of the column alias
- The second heading appears as PROFIT
- If the column alias was not used how would the output appear?



# Using Arithmetic Operations



The screenshot shows a database query tool interface. At the top, there is a toolbar with various icons. Below it, a tab labeled 'Worksheet' is active, showing a SQL query in a text editor. The query is:

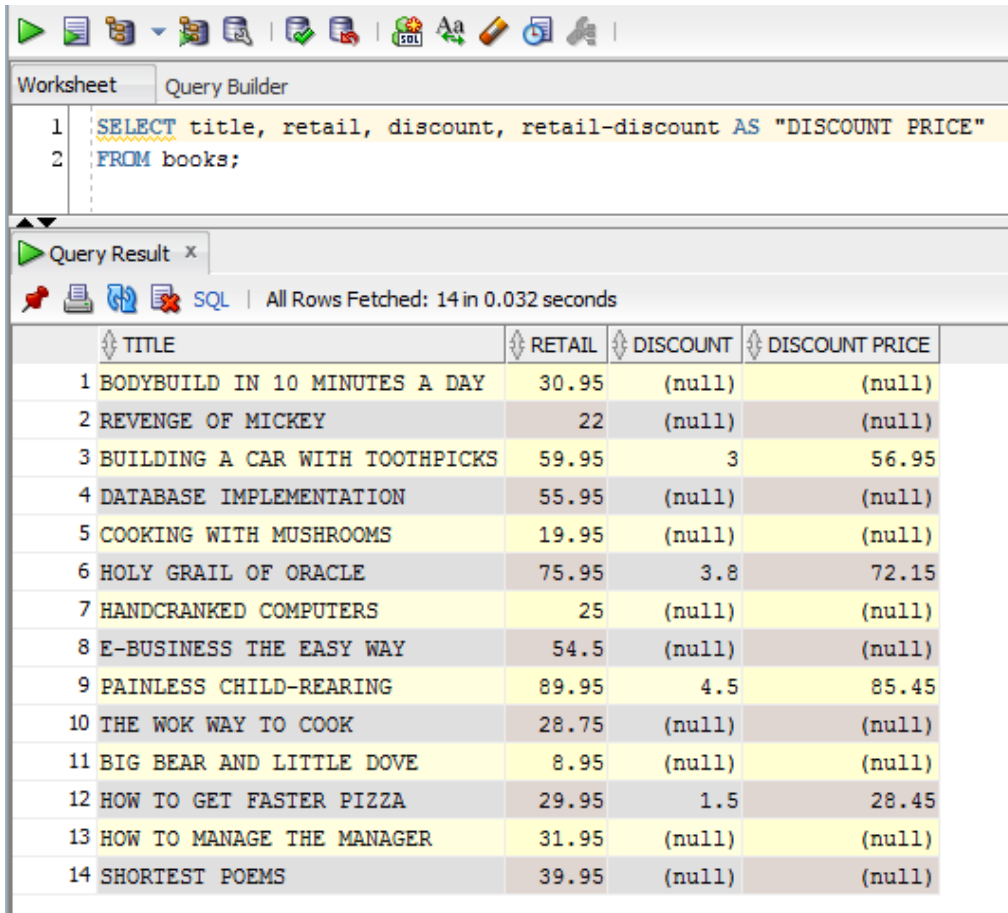
```
1 SELECT title, retail, discount, retail-discount
2 FROM books;
```

Below the query editor, there is a 'Query Result' tab. It shows the results of the query in a table. The table has four columns: TITLE, RETAIL, DISCOUNT, and RETAIL-DISCOUNT. The results are as follows:

	TITLE	RETAIL	DISCOUNT	RETAIL-DISCOUNT
1	BODYBUILD IN 10 MINUTES A DAY	30.95	(null)	(null)
2	REVENGE OF MICKEY	22	(null)	(null)
3	BUILDING A CAR WITH TOOTHPICKS	59.95	3	56.95
4	DATABASE IMPLEMENTATION	55.95	(null)	(null)
5	COOKING WITH MUSHROOMS	19.95	(null)	(null)
6	HOLY GRAIL OF ORACLE	75.95	3.8	72.15
7	HANDCRANKED COMPUTERS	25	(null)	(null)
8	E-BUSINESS THE EASY WAY	54.5	(null)	(null)
9	PAINLESS CHILD-REARING	89.95	4.5	85.45
10	THE WOK WAY TO COOK	28.75	(null)	(null)
11	BIG BEAR AND LITTLE DOVE	8.95	(null)	(null)
12	HOW TO GET FASTER PIZZA	29.95	1.5	28.45
13	HOW TO MANAGE THE MANAGER	31.95	(null)	(null)
14	SHORTEST POEMS	39.95	(null)	(null)

- The heading shows what appears in the Select CLAUSE, RETAIL-DISCOUNT
- An alias would help to make the result set more understandable
- Which is better to see an alias of DISCOUNT PRICE or RETAIL-DISCOUNT
- DISCOUNT PRICE shows a clearer picture of the data in the column
- Alias values are used quite frequently when arithmetic statements are performed

# Using Arithmetic Operations



The screenshot shows a database query builder interface. The top toolbar contains various icons for file operations, execution, and formatting. Below the toolbar, the 'Worksheet' tab is active, displaying a SQL query in a text editor. The query is as follows:

```
1 SELECT title, retail, discount, retail-discount AS "DISCOUNT PRICE"
2 FROM books;
```

Below the query editor, the 'Query Result' tab is active, showing the results of the query. The results are displayed in a table with 4 columns: TITLE, RETAIL, DISCOUNT, and DISCOUNT PRICE. The table contains 14 rows of data. The 'DISCOUNT PRICE' column is calculated as 'retail-discount'.

	TITLE	RETAIL	DISCOUNT	DISCOUNT PRICE
1	BODYBUILD IN 10 MINUTES A DAY	30.95	(null)	(null)
2	REVENGE OF MICKEY	22	(null)	(null)
3	BUILDING A CAR WITH TOOTHPICKS	59.95	3	56.95
4	DATABASE IMPLEMENTATION	55.95	(null)	(null)
5	COOKING WITH MUSHROOMS	19.95	(null)	(null)
6	HOLY GRAIL OF ORACLE	75.95	3.8	72.15
7	HANDCRANKED COMPUTERS	25	(null)	(null)
8	E-BUSINESS THE EASY WAY	54.5	(null)	(null)
9	PAINLESS CHILD-REARING	89.95	4.5	85.45
10	THE WOK WAY TO COOK	28.75	(null)	(null)
11	BIG BEAR AND LITTLE DOVE	8.95	(null)	(null)
12	HOW TO GET FASTER PIZZA	29.95	1.5	28.45
13	HOW TO MANAGE THE MANAGER	31.95	(null)	(null)
14	SHORTEST POEMS	39.95	(null)	(null)

- Alias being used to clarify result in column

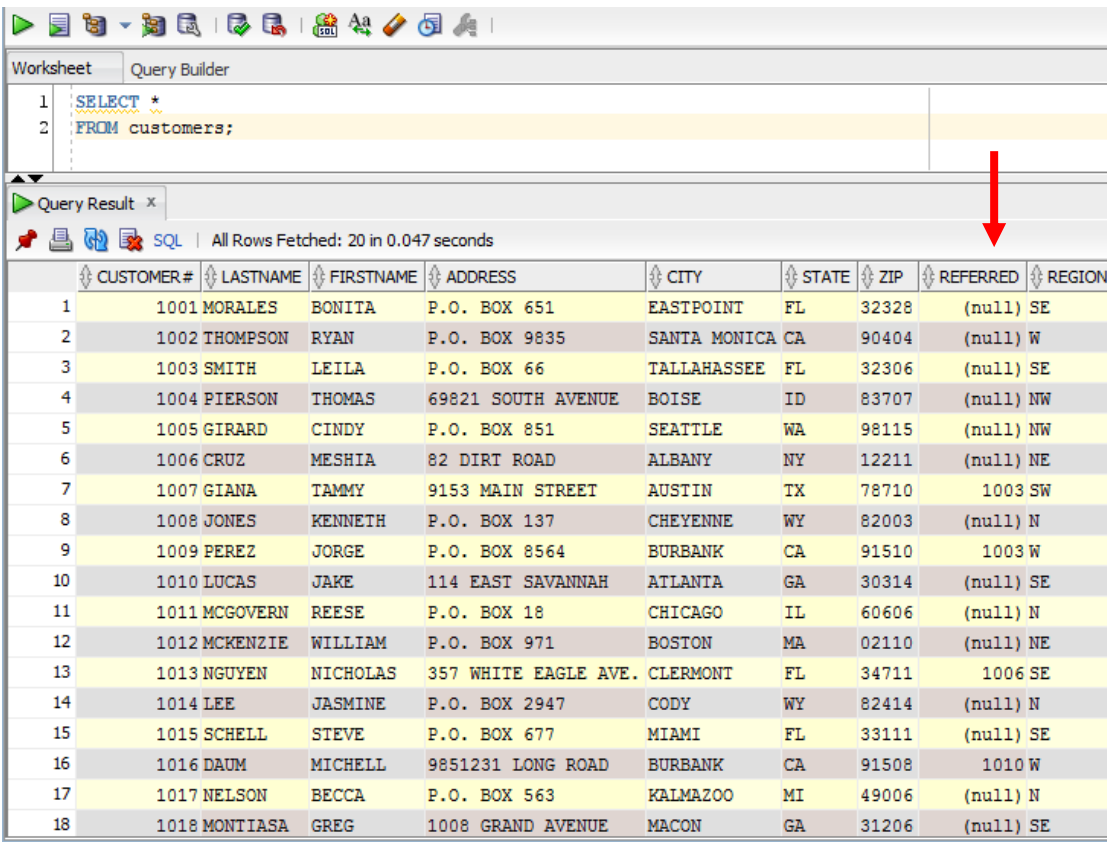
# Order of Operations

- Moving from left to right in the arithmetic equation, any required multiplication and division operations are solved first
- Addition and subtraction operations are solved after multiplication and division, again moving from left to right in the equation
- To override this order of operations, you use parentheses to enclose a portion that should be calculated first

# NULL Values

- If no value is entered for a column in a row of data the value is considered NULL
- It indicates the absence of data
- In the next slide you will notice in the COMM column the word (null) for any row value that has a NULL value for the COMM

# NULL Values



Worksheet Query Builder

```
1 SELECT *
2 FROM customers;
```

Query Result x

SQL | All Rows Fetched: 20 in 0.047 seconds

	CUSTOMER#	LASTNAME	FIRSTNAME	ADDRESS	CITY	STATE	ZIP	REFERRED	REGION
1	1001	MORALES	BONITA	P.O. BOX 651	EASTPOINT	FL	32328	(null)	SE
2	1002	THOMPSON	RYAN	P.O. BOX 9835	SANTA MONICA	CA	90404	(null)	W
3	1003	SMITH	LEILA	P.O. BOX 66	TALLAHASSEE	FL	32306	(null)	SE
4	1004	PIERSON	THOMAS	69821 SOUTH AVENUE	BOISE	ID	83707	(null)	NW
5	1005	GIRARD	CINDY	P.O. BOX 851	SEATTLE	WA	98115	(null)	NW
6	1006	CRUZ	MESHIA	82 DIRT ROAD	ALBANY	NY	12211	(null)	NE
7	1007	GIANA	TAMMY	9153 MAIN STREET	AUSTIN	TX	78710	1003	SW
8	1008	JONES	KENNETH	P.O. BOX 137	CHEYENNE	WY	82003	(null)	N
9	1009	PEREZ	JORGE	P.O. BOX 8564	BURBANK	CA	91510	1003	W
10	1010	LUCAS	JAKE	114 EAST SAVANNAH	ATLANTA	GA	30314	(null)	SE
11	1011	MCGOVERN	REESE	P.O. BOX 18	CHICAGO	IL	60606	(null)	N
12	1012	MCKENZIE	WILLIAM	P.O. BOX 971	BOSTON	MA	02110	(null)	NE
13	1013	NGUYEN	NICHOLAS	357 WHITE EAGLE AVE.	CLERMONT	FL	34711	1006	SE
14	1014	LEE	JASMINE	P.O. BOX 2947	CODY	WY	82414	(null)	N
15	1015	SHELL	STEVE	P.O. BOX 677	MIAMI	FL	33111	(null)	SE
16	1016	DAUM	MICHELL	9851231 LONG ROAD	BURBANK	CA	91508	1010	W
17	1017	NELSON	BECCA	P.O. BOX 563	KALMAZOO	MI	49006	(null)	N
18	1018	MONTIASA	GREG	1008 GRAND AVENUE	MACON	GA	31206	(null)	SE

Notice the word NULL for any value that is not defined

NULL is not stored in the database, it appears for our use to give a visual indication that the column contains a NULL value

# NULL Values

Worksheet

Query Builder

1

SELECT \*

2

FROM customers;

Query Result x

Script Output x

CUSTOMER#	LASTNAME	FIRSTNAME	ADDRESS	CITY	ST	ZIP	REFERRED	RE
1001	MORALES	BONITA	P.O. BOX 651	EASTPOINT	FL	32328		SE
1002	THOMPSON	RYAN	P.O. BOX 9835	SANTA MONICA	CA	90404		W
1003	SMITH	LEILA	P.O. BOX 66	TALLAHASSEE	FL	32306		SE
1004	PIERSON	THOMAS	69821 SOUTH AVENUE	BOISE	ID	83707		NW
1005	GIRARD	CINDY	P.O. BOX 851	SEATTLE	WA	98115		NW
1006	CRUZ	MESHIA	82 DIRT ROAD	ALBANY	NY	12211		NE
1007	GIANA	TAMMY	9153 MAIN STREET	AUSTIN	TX	78710	1003	SW
1008	JONES	KENNETH	P.O. BOX 137	CHEYENNE	WY	82003		N
1009	PEREZ	JORGE	P.O. BOX 8564	BURBANK	CA	91510	1003	W
1010	LUCAS	JAKE	114 EAST SAVANNAH	ATLANTA	GA	30314		SE
1011	MCGOVERN	REESE	P.O. BOX 18	CHICAGO	IL	60606		N

CUSTOMER#	LASTNAME	FIRSTNAME	ADDRESS	CITY	ST	ZIP	REFERRED	RE
1012	MCKENZIE	WILLIAM	P.O. BOX 971	BOSTON	MA	02110		NE
1013	NGUYEN	NICHOLAS	357 WHITE EAGLE AVE.	CLERMONT	FL	34711	1006	SE
1014	LEE	JASMINE	P.O. BOX 2947	CODY	WY	82414		N
1015	SCHELL	STEVE	P.O. BOX 677	MIAMI	FL	33111		SE
1016	DAUM	MICHELL	9851231 LONG ROAD	BURBANK	CA	91508	1010	W
1017	NELSON	BECCA	P.O. BOX 563	KALMAZOO	MI	49006		N

- If I use the alternate method to execute a query, notice the NULL values appear with no visible value or any indication the column is NULL

# NULL Values

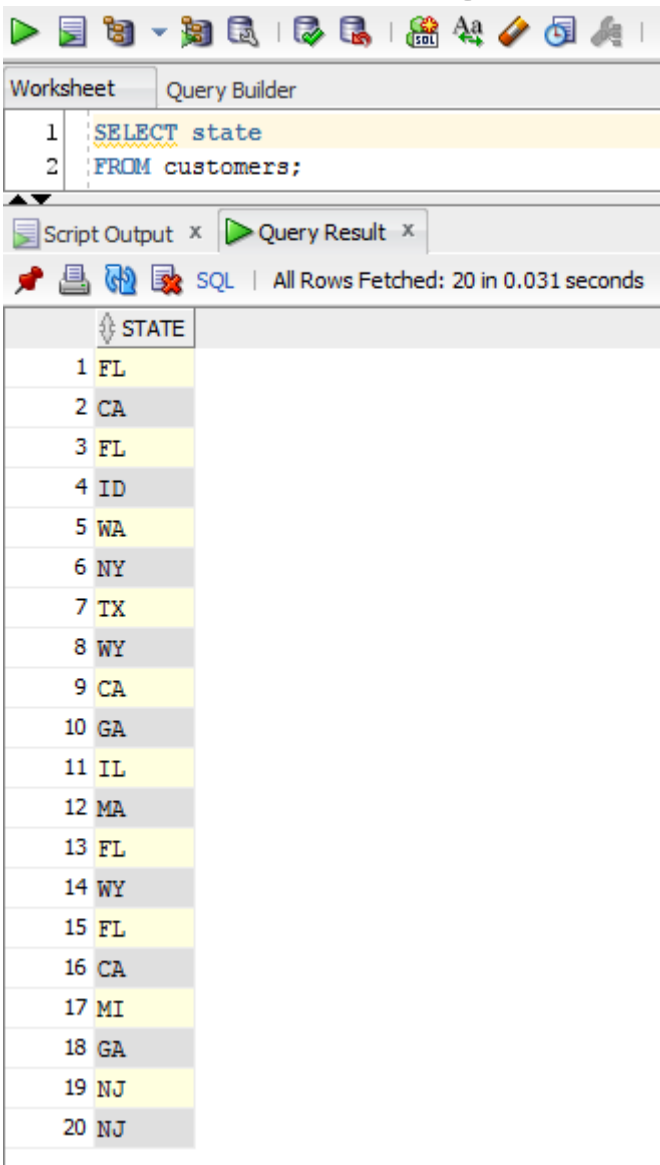
- NULL values can lead to undesirable results in operations
- For example, what if you need to an employee's total salary, as we saw this was difficult to determine
- Seems simple enough, but check out the result for every row where the REFERRED value was NULL
- This can be fixed, we shall see later on a function can be used to allow you to substitute a value for a NULL

# Using DISTINCT and UNIQUE

- Suppose you want to know which are the various jobs the employees can have
- You want to list only the jobs, not the employee information. This could be done by listing only the JOB column of the EMPLOYEES table
- This is done on the next slide. Do you notice a problem?



# Using DISTINCT and UNIQUE

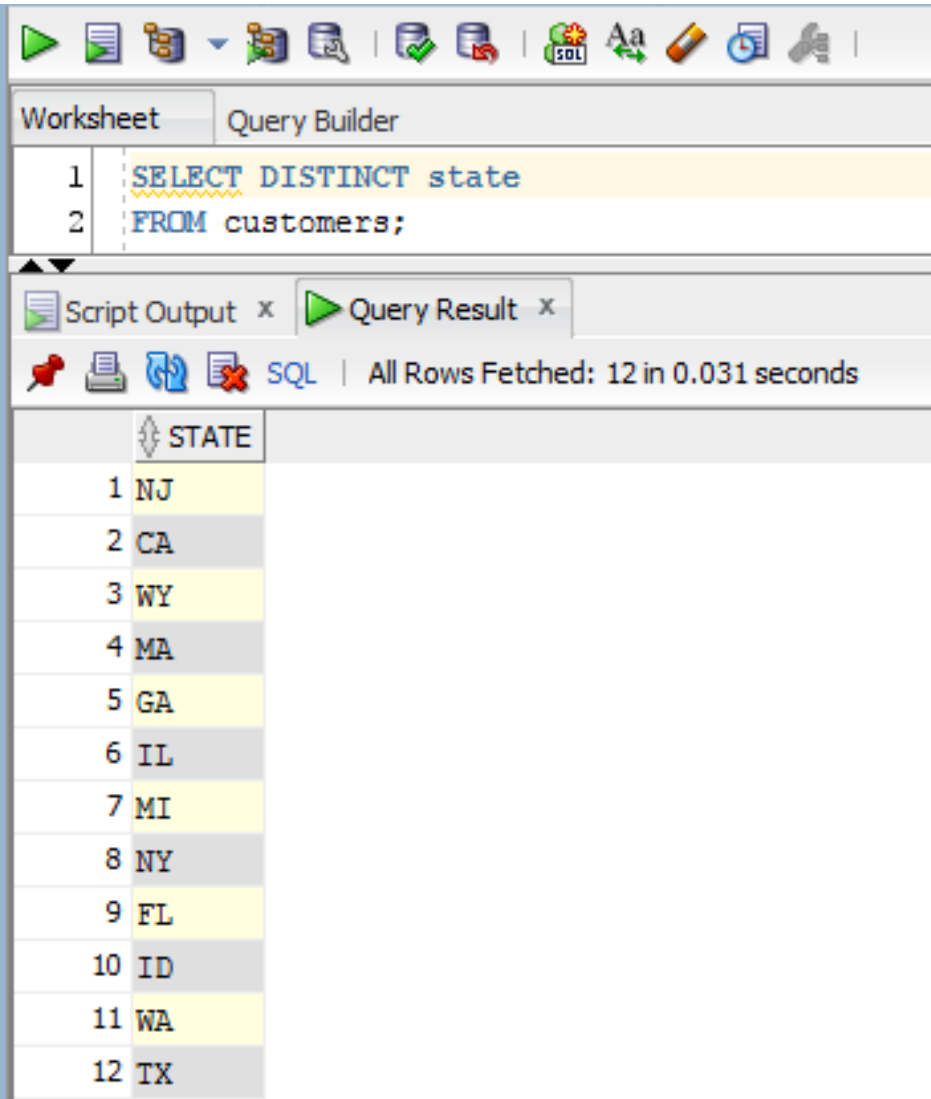


The screenshot shows a SQL query editor with a toolbar at the top. Below the toolbar, there are tabs for 'Worksheet' and 'Query Builder'. The 'Query Builder' tab is active, showing a query with two lines: 'SELECT state' and 'FROM customers;'. Below the query, there are tabs for 'Script Output' and 'Query Result'. The 'Query Result' tab is active, showing a table with 20 rows. The table has a single column labeled 'STATE'. The rows contain the following state abbreviations: FL, CA, FL, ID, WA, NY, TX, WY, CA, GA, IL, MA, FL, WY, FL, CA, MI, GA, NJ, NJ.

	STATE
1	FL
2	CA
3	FL
4	ID
5	WA
6	NY
7	TX
8	WY
9	CA
10	GA
11	IL
12	MA
13	FL
14	WY
15	FL
16	CA
17	MI
18	GA
19	NJ
20	NJ

- The output shows one row of output for each row in the table
- 20 rows shows 20 state values, one for each row
- If the question was asked, show the jobs that my employees may have, would I want to see the value for each employee or would I only need to see individual state values?
- I am glad I only have 14 customers showing and not 50,000 customers
- How can I modify the query to show each state only once?

# Using DISTINCT and UNIQUE



The screenshot shows a database query builder interface. At the top, there is a toolbar with various icons. Below the toolbar, there are two tabs: 'Worksheet' and 'Query Builder'. The 'Query Builder' tab is active, and it contains a SQL query editor with the following text:

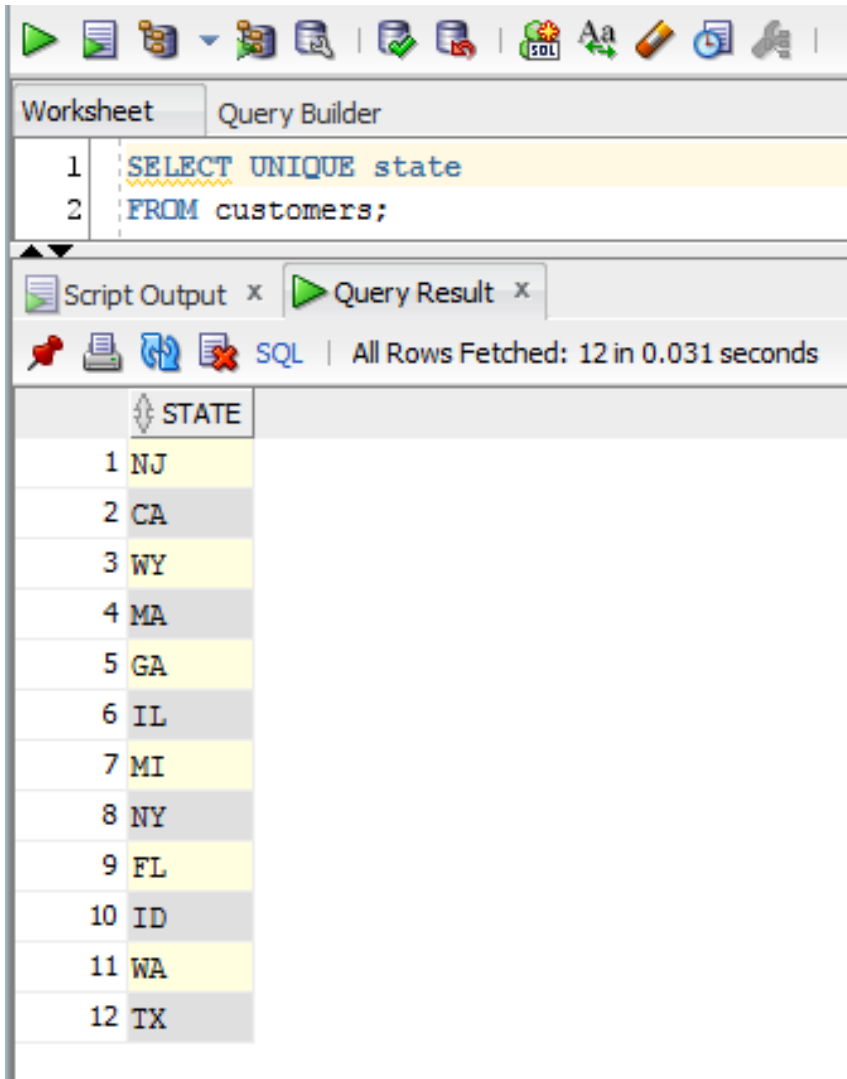
```
1 SELECT DISTINCT state
2 FROM customers;
```

Below the query editor, there is a 'Script Output' tab and a 'Query Result' tab. The 'Query Result' tab is active, and it displays the results of the query. The results are shown in a table with a single column labeled 'STATE'. The table contains 12 rows, each with a unique state abbreviation. The states are: NJ, CA, WY, MA, GA, IL, MI, NY, FL, ID, WA, and TX. The table is displayed with alternating yellow and white rows.

	STATE
1	NJ
2	CA
3	WY
4	MA
5	GA
6	IL
7	MI
8	NY
9	FL
10	ID
11	WA
12	TX

- In this output I only see an individual row for each state
- There are no duplicate values displayed for the state column
- Even if I had 50,000 customers in the table, and all customers were in one of the 12 different states showing then I would only see 12 rows of output still
- Any duplicate values are suppressed

# Using DISTINCT and UNIQUE



The screenshot shows a database query tool interface. At the top is a toolbar with various icons. Below it is a tabbed interface with 'Worksheet' and 'Query Builder' tabs. The 'Query Builder' tab is active, showing a SQL query in a text area:

```
1 SELECT UNIQUE state
2 FROM customers;
```

Below the query area is a 'Script Output' tab and a 'Query Result' tab. The 'Query Result' tab is active, showing a table with 12 rows. The table has a single column labeled 'STATE'. The rows contain the following state abbreviations: NJ, CA, WY, MA, GA, IL, MI, NY, FL, ID, WA, and TX. The first row is highlighted in yellow.

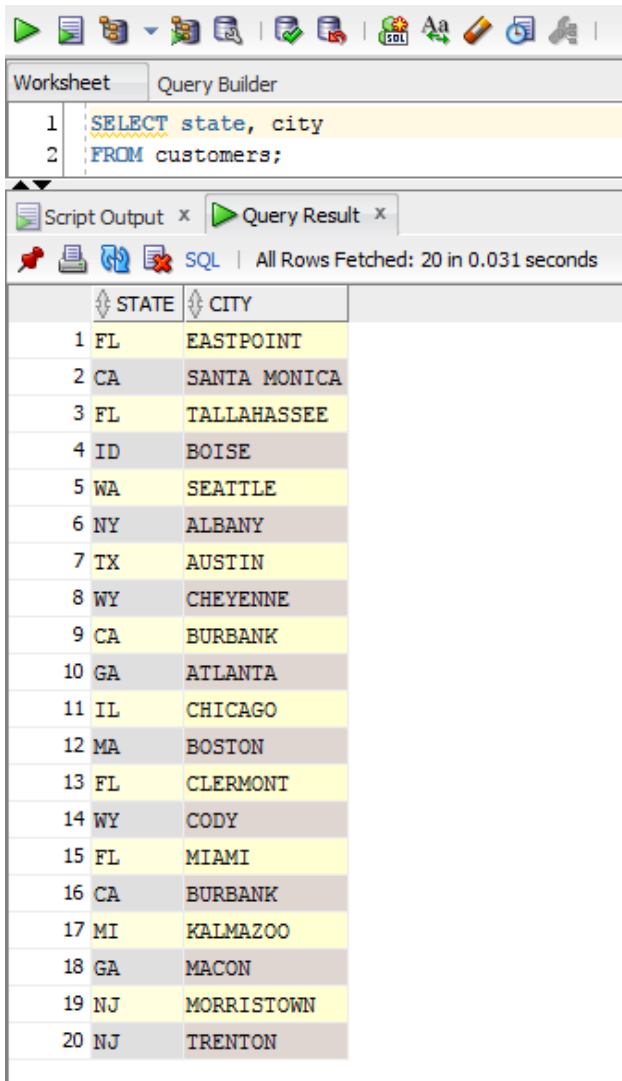
	STATE
1	NJ
2	CA
3	WY
4	MA
5	GA
6	IL
7	MI
8	NY
9	FL
10	ID
11	WA
12	TX

- The UNIQUE keyword displays the same output as the DISTINCT keyword
- Both perform the same function and are interchangeable

# Using DISTINCT and UNIQUE

- The **DISTINCT** keyword is applied to **all** columns listed in the **SELECT** statement
- So, if multiple columns were used in the **SELECT**, **each** possible result would be tested for uniqueness. Only the ones that are unique for **all** columns will be returned

# Using DISTINCT

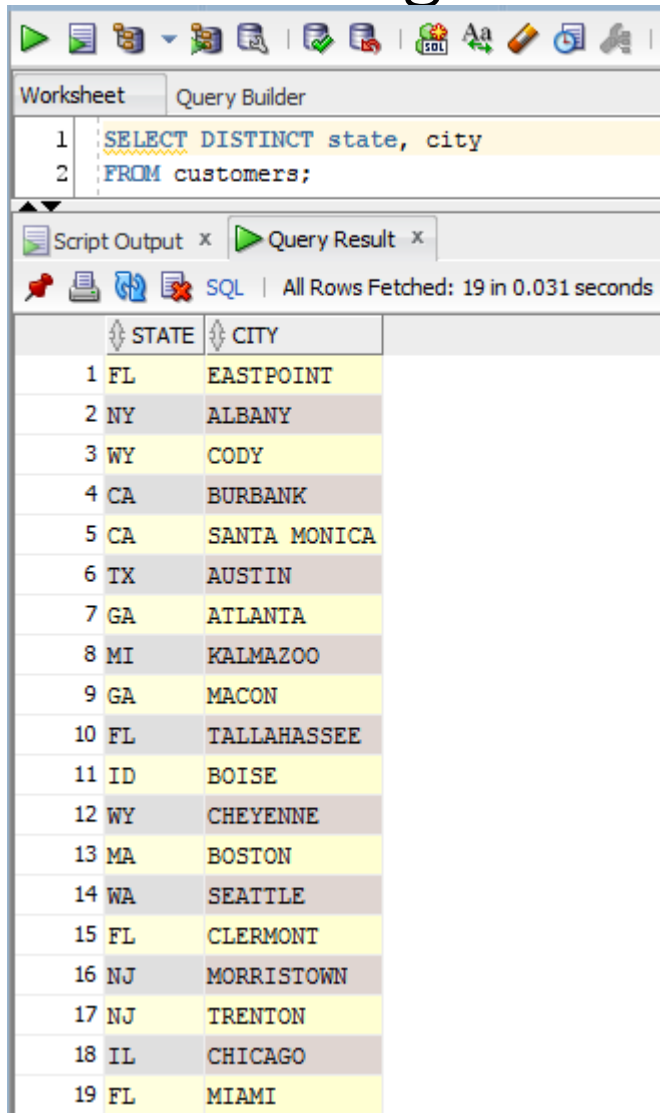


The screenshot shows a database query tool interface. At the top is a toolbar with various icons. Below it is a tabbed interface with 'Worksheet' and 'Query Builder' tabs. The 'Query Builder' tab contains a SQL query: `SELECT state, city` on line 1 and `FROM customers;` on line 2. Below the query is a 'Script Output' tab and a 'Query Result' tab. The 'Query Result' tab shows the results of the query, with a status bar indicating 'All Rows Fetched: 20 in 0.031 seconds'. The results are displayed in a table with two columns: 'STATE' and 'CITY'. The table contains 20 rows of data, each with a row number, a state abbreviation, and a city name.

	STATE	CITY
1	FL	EASTPOINT
2	CA	SANTA MONICA
3	FL	TALLAHASSEE
4	ID	BOISE
5	WA	SEATTLE
6	NY	ALBANY
7	TX	AUSTIN
8	WY	CHEYENNE
9	CA	BURBANK
10	GA	ATLANTA
11	IL	CHICAGO
12	MA	BOSTON
13	FL	CLERMONT
14	WY	CODY
15	FL	MIAMI
16	CA	BURBANK
17	MI	KALMAZOO
18	GA	MACON
19	NJ	MORRISTOWN
20	NJ	TRENTON

Here I am showing the STATE and CITY values, there is one row for every CUSTOMER

# Using DISTINCT and UNIQUE



The screenshot shows a SQL query builder window. The query is: `SELECT DISTINCT state, city FROM customers;`. The results are displayed in a table with columns STATE and CITY. The results are 19 rows long, showing unique combinations of state and city.

	STATE	CITY
1	FL	EASTPOINT
2	NY	ALBANY
3	WY	CODY
4	CA	BURBANK
5	CA	SANTA MONICA
6	TX	AUSTIN
7	GA	ATLANTA
8	MI	KALMAZOO
9	GA	MACON
10	FL	TALLAHASSEE
11	ID	BOISE
12	WY	CHEYENNE
13	MA	BOSTON
14	WA	SEATTLE
15	FL	CLERMONT
16	NJ	MORRISTOWN
17	NJ	TRENTON
18	IL	CHICAGO
19	FL	MIAMI

- Instead of the 20 rows for all customers or the 12 rows for the 12 different states shown before or the unique results of CITY and STATE combinations
- This tells me that two of my customers have the same STATE and CITY
- I do not know which ones are the same, only that two must be the same

# DISTINCT versus UNIQUE

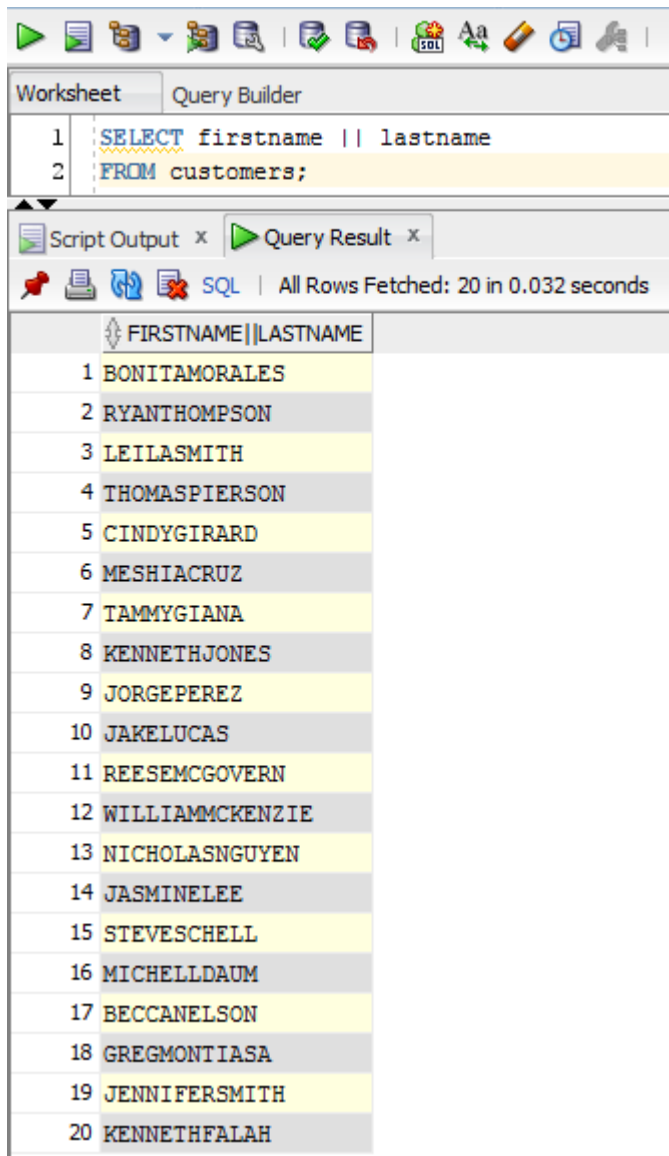
- I was asked, “what is the difference between DISTINCT and UNIQUE?”
- UNIQUE was Oracle’s traditional method for removing duplicate values, and only Oracle uses this method
- DISTINCT was added when Oracle adopted the ANSI specifications
- So DISTINCT will work in other SQL versions, where UNIQUE will only function with Oracle

# Concatenation

- So far, each field has been placed in a column of its own
- In certain situations, you might want to display the contents of each field so they appear right next to each other, separated by a single space or comma
- Combining columns is called **concatenation**
- The concatenation operator in Oracle is two vertical bars **||**



# Concatenation

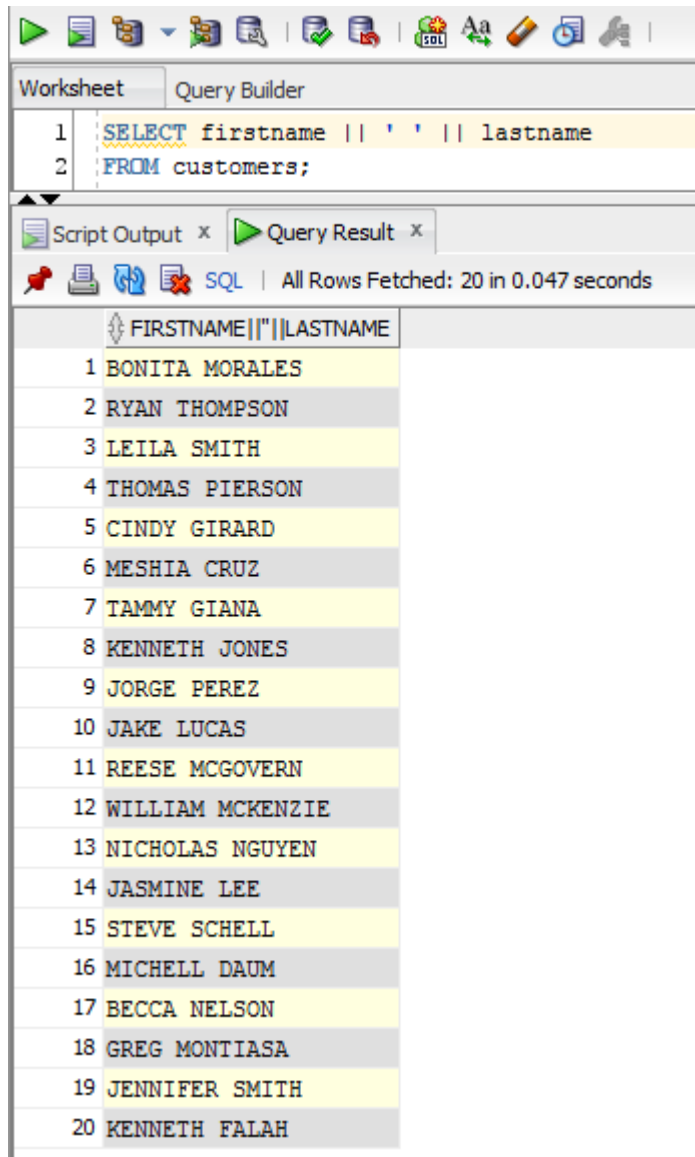


The screenshot shows a database query tool interface. At the top, there's a toolbar with various icons. Below it, a tab labeled 'Query Builder' is active. The query editor shows two lines of SQL code: `1 SELECT firstname || lastname` and `2 FROM customers;`. Below the query editor, there's a 'Script Output' tab and a 'Query Result' tab. The 'Query Result' tab is active, showing a table with 20 rows. The table has two columns: 'FIRSTNAME' and 'LASTNAME'. The results are concatenated as 'FIRSTNAME||LASTNAME'. The rows are numbered 1 through 20.

	FIRSTNAME  LASTNAME
1	BONITAMORALES
2	RYANTHOMPSON
3	LEILASMITH
4	THOMASPIERSON
5	CINDYGIRARD
6	MESHIACRUZ
7	TAMMYGIANA
8	KENNETHJONES
9	JORGEPEREZ
10	JAKELUCAS
11	REESEMCGOVERN
12	WILLIAMMCKENZIE
13	NICHOLASNGUYEN
14	JASMINELEE
15	STEVESCHELL
16	MICHELLEDAUM
17	BECCANELSON
18	GREGMONTIASA
19	JENNIFERSMITH
20	KENNETHFALAH

- This did as I requested, it concatenated the first name and the last name together
- It is not easy to read, can it be improved?
- There should be a space between the two values
- Time to try again to improve the readability of our result

# Concatenation



The screenshot shows a database query tool interface. At the top, there's a toolbar with various icons. Below it, a tab labeled 'Query Builder' is active. The query editor shows two lines of SQL code: `1 SELECT firstname || ' ' || lastname` and `2 FROM customers;`. Below the query editor, there's a 'Script Output' and 'Query Result' tab. The 'Query Result' tab is active, showing a table with one column named 'FIRSTNAME||" "||LASTNAME'. The table contains 20 rows of data, each with a row number and a concatenated string of a first name, a space, and a last name. The status bar at the bottom indicates 'All Rows Fetched: 20 in 0.047 seconds'.

	FIRSTNAME  " "  LASTNAME
1	BONITA MORALES
2	RYAN THOMPSON
3	LEILA SMITH
4	THOMAS PIERSON
5	CINDY GIRARD
6	MESHIA CRUZ
7	TAMMY GIANA
8	KENNETH JONES
9	JORGE PEREZ
10	JAKE LUCAS
11	REESE MCGOVERN
12	WILLIAM MCKENZIE
13	NICHOLAS NGUYEN
14	JASMINE LEE
15	STEVE SCHELL
16	MICHELL DAUM
17	BECCA NELSON
18	GREG MONTIASA
19	JENNIFER SMITH
20	KENNETH FALAH

- This looks better, I have the space between two values
- Could it still be improved?
- What about the headings?
- How can I fix the heading?

# Concatenation

- The previous slide took two concatenation operations
- In the previous slide, a **string literal** was inserted into the output to put a blank space between the two fields
- All string literals are enclosed in **single quotation marks**, the single quotes have a space inside so I end up with a space in my result
- On the previous slide, the lastname is concatenated to the blank space, then the blank space is concatenated to the firstname

# Concatenation

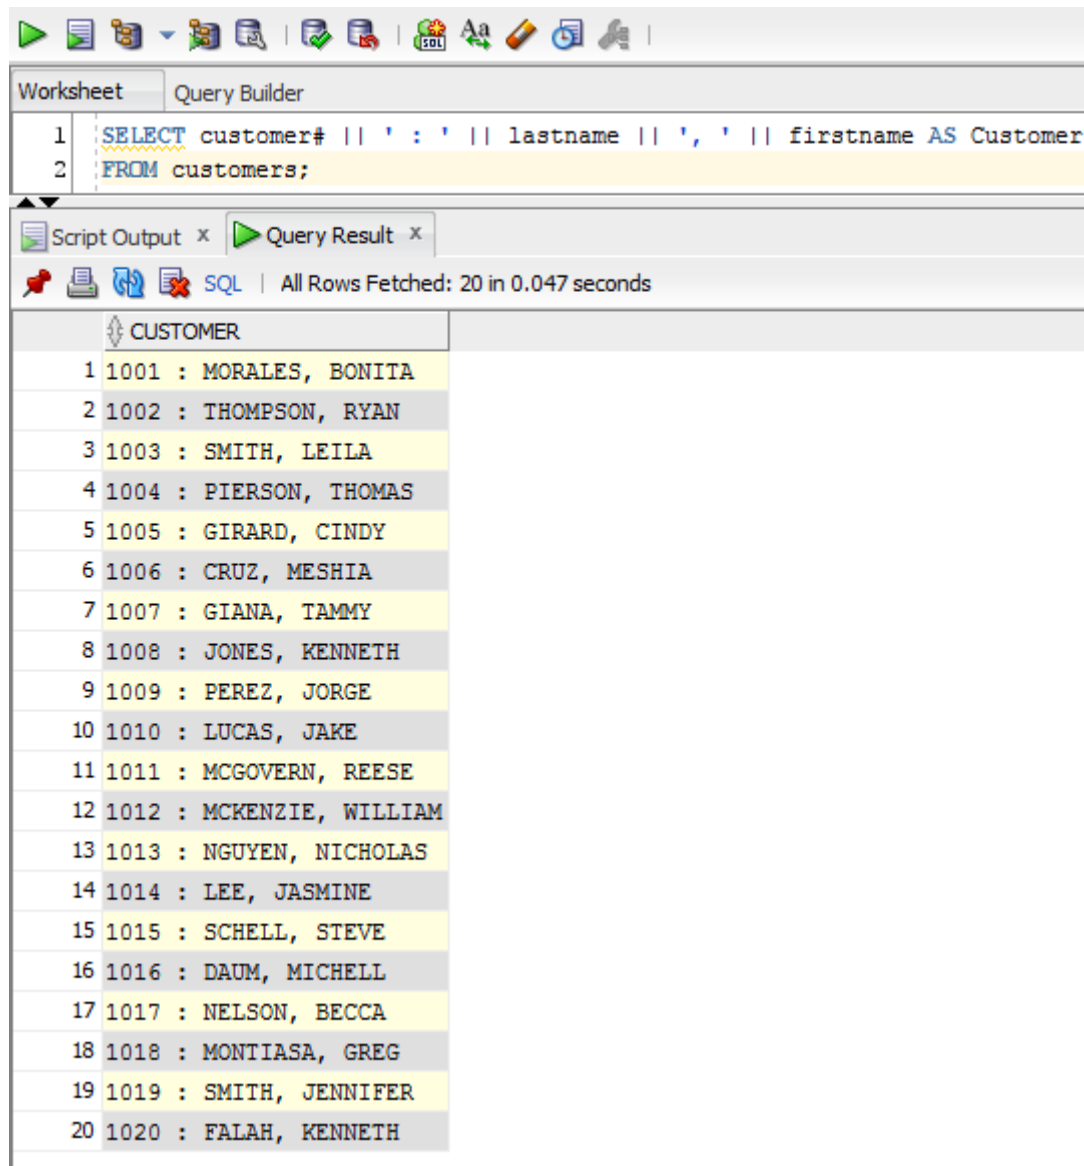


The screenshot shows a database query tool interface. At the top, there is a toolbar with various icons. Below it, a tab labeled 'Query Builder' is active. The query editor shows two lines of SQL code: `1 SELECT firstname || ' ' || lastname AS "Full Name"` and `2 FROM customers;`. Below the query editor, there are tabs for 'Script Output' and 'Query Result'. The 'Query Result' tab is selected, showing a table with one column named 'Full Name'. The table contains 20 rows of data, each with a row number and a full name. The names are: 1 BONITA MORALES, 2 RYAN THOMPSON, 3 LEILA SMITH, 4 THOMAS PIERSON, 5 CINDY GIRARD, 6 MESHIA CRUZ, 7 TAMMY GIANA, 8 KENNETH JONES, 9 JORGE PEREZ, 10 JAKE LUCAS, 11 REESE MCGOVERN, 12 WILLIAM MCKENZIE, 13 NICHOLAS NGUYEN, 14 JASMINE LEE, 15 STEVE SCHELL, 16 MICHELL DAUM, 17 BECCA NELSON, 18 GREG MONTIASA, 19 JENNIFER SMITH, and 20 KENNETH FALAH.

	Full Name
1	BONITA MORALES
2	RYAN THOMPSON
3	LEILA SMITH
4	THOMAS PIERSON
5	CINDY GIRARD
6	MESHIA CRUZ
7	TAMMY GIANA
8	KENNETH JONES
9	JORGE PEREZ
10	JAKE LUCAS
11	REESE MCGOVERN
12	WILLIAM MCKENZIE
13	NICHOLAS NGUYEN
14	JASMINE LEE
15	STEVE SCHELL
16	MICHELL DAUM
17	BECCA NELSON
18	GREG MONTIASA
19	JENNIFER SMITH
20	KENNETH FALAH

- Here an alias is used to make the output more readable
- It is always advisable to use column aliases to replace headings for expressions or concatenated values that could appear in the headings when concatenation is used

# Concatenation

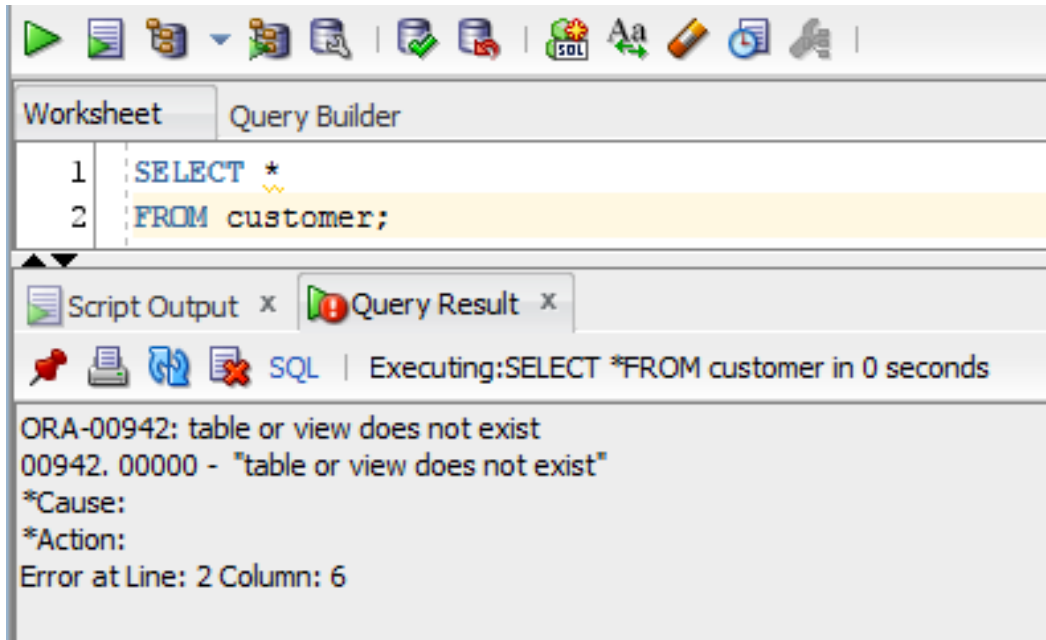


The screenshot shows a SQL query builder interface. The 'Query Builder' tab is active, displaying a query with two lines: `1 SELECT customer# || ' : ' || lastname || ', ' || firstname AS Customer` and `2 FROM customers;`. Below the query, the 'Query Result' tab shows the results of the query. The results are displayed in a table with a single column labeled 'CUSTOMER'. The table contains 20 rows of data, each showing a customer ID followed by a colon, a space, the last name, a comma, a space, and the first name.

	CUSTOMER
1	1001 : MORALES, BONITA
2	1002 : THOMPSON, RYAN
3	1003 : SMITH, LEILA
4	1004 : PIERSON, THOMAS
5	1005 : GIRARD, CINDY
6	1006 : CRUZ, MESHIA
7	1007 : GIANA, TAMMY
8	1008 : JONES, KENNETH
9	1009 : PEREZ, JORGE
10	1010 : LUCAS, JAKE
11	1011 : MCGOVERN, REESE
12	1012 : MCKENZIE, WILLIAM
13	1013 : NGUYEN, NICHOLAS
14	1014 : LEE, JASMINE
15	1015 : SCHELL, STEVE
16	1016 : DAUM, MICHELL
17	1017 : NELSON, BECCA
18	1018 : MONTIASA, GREG
19	1019 : SMITH, JENNIFER
20	1020 : FALAH, KENNETH

- This is another variation where string literals have been placed into the result

# Syntax Errors



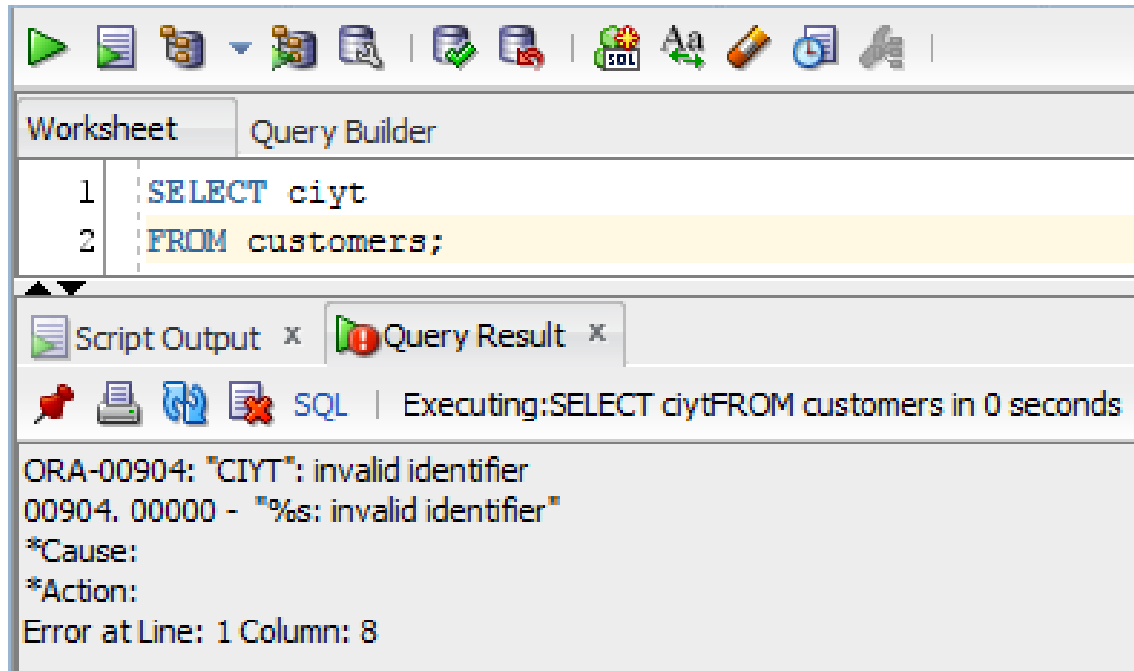
The query to the right produces an error ORA-00942 table or view does not exist

Table name in this case is incorrect it is CUSTOMERS not CUSTOMER

Correct and the re-execute the command

Check your tree on the left to see the table names

# Syntax Errors



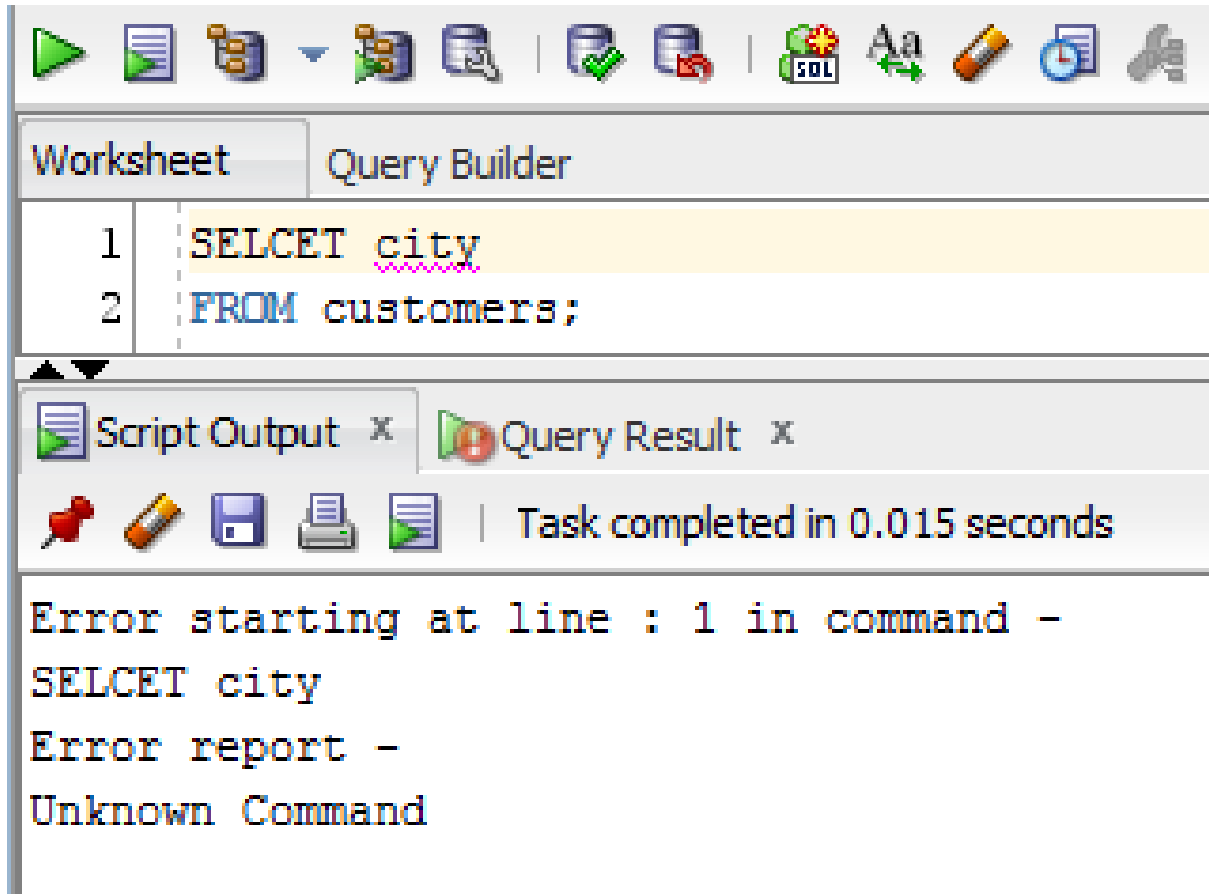
The query to the right has produced an error  
ORA-00904: CIYTinvalid identifier

Column name is misspelled, should be CITY

Again check the tree on the left to see the column names in the table

Could also do a DESCRIBE customers; to see correct columns

# Syntax Errors



The keyword SELECT is misspelled

Gives error as Unknown Command

Notice the keyword did turn BLUE indicating the error, you will notice all keywords will display in BLUE to indicate they are keywords

Fix and re-execute

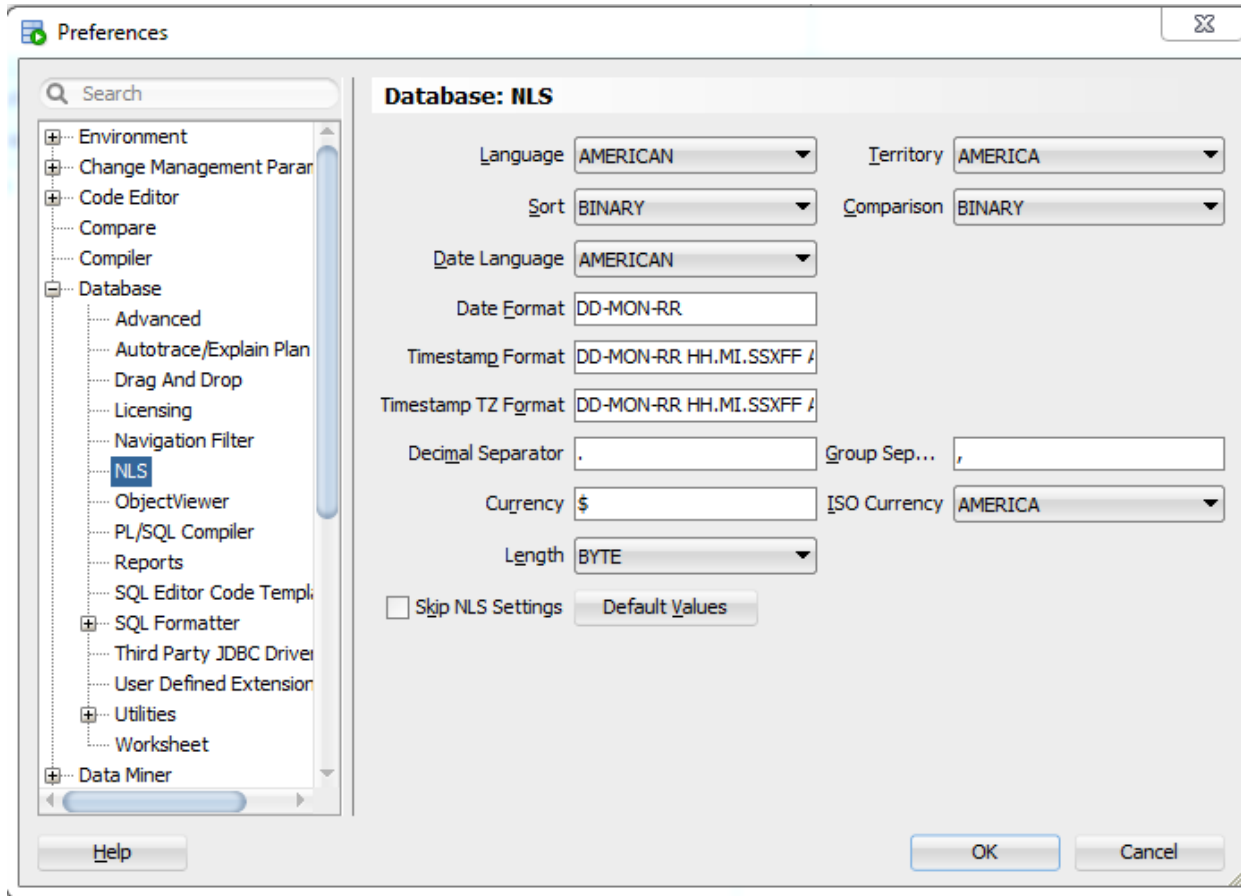
Notice the red squiggle after the misspelled keyword



# Settings

- It is important to make sure your date values are set correctly
- The default date format Oracle uses is DD-MON-RR
- When you install SQL Developer it looks at the Windows settings and sets the date format according to country
- ENGLISH CANADA is a different value
- Go to the Tools menu, select PREFERENCES, expand DATABASE then select NLS
- This is the important setting to check, please make sure otherwise you may get errors in your outputs involving date values

# Settings



Set the following:  
Date Format  
Timestamp Format  
Timestamp TZ  
Format

They should all  
have DD-MON-RR

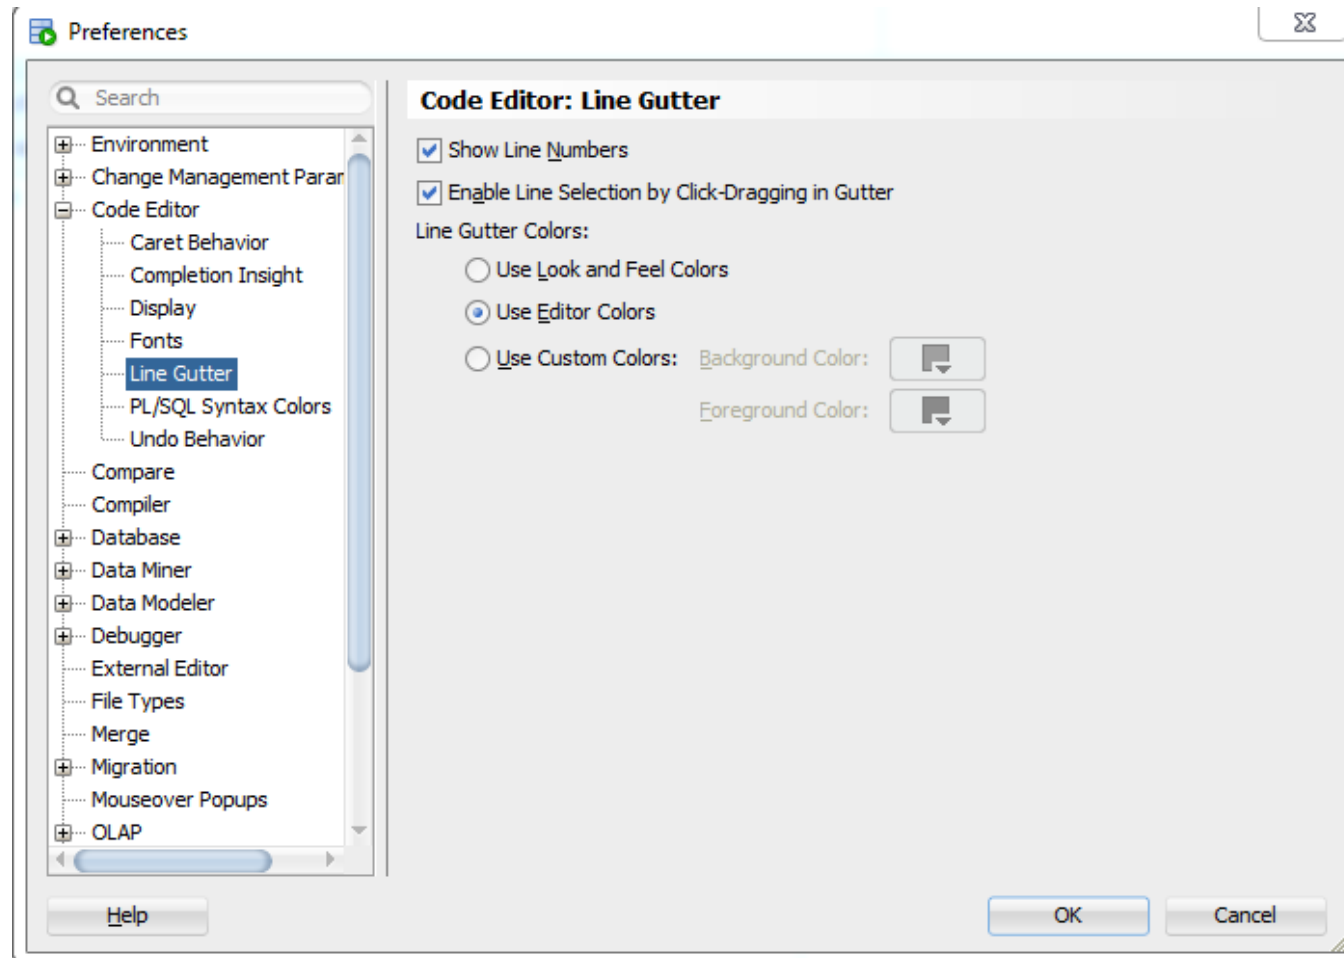
If yours is different  
please modify then  
click OK

On the Timestamps  
you only need to  
modify the left  
most values

# Line Numbers

- While you are in the Preferences click on the Code Editor option
- Then select Line Gutter
- Check the Show Line Numbers check box
- This will then show line numbers in the editor

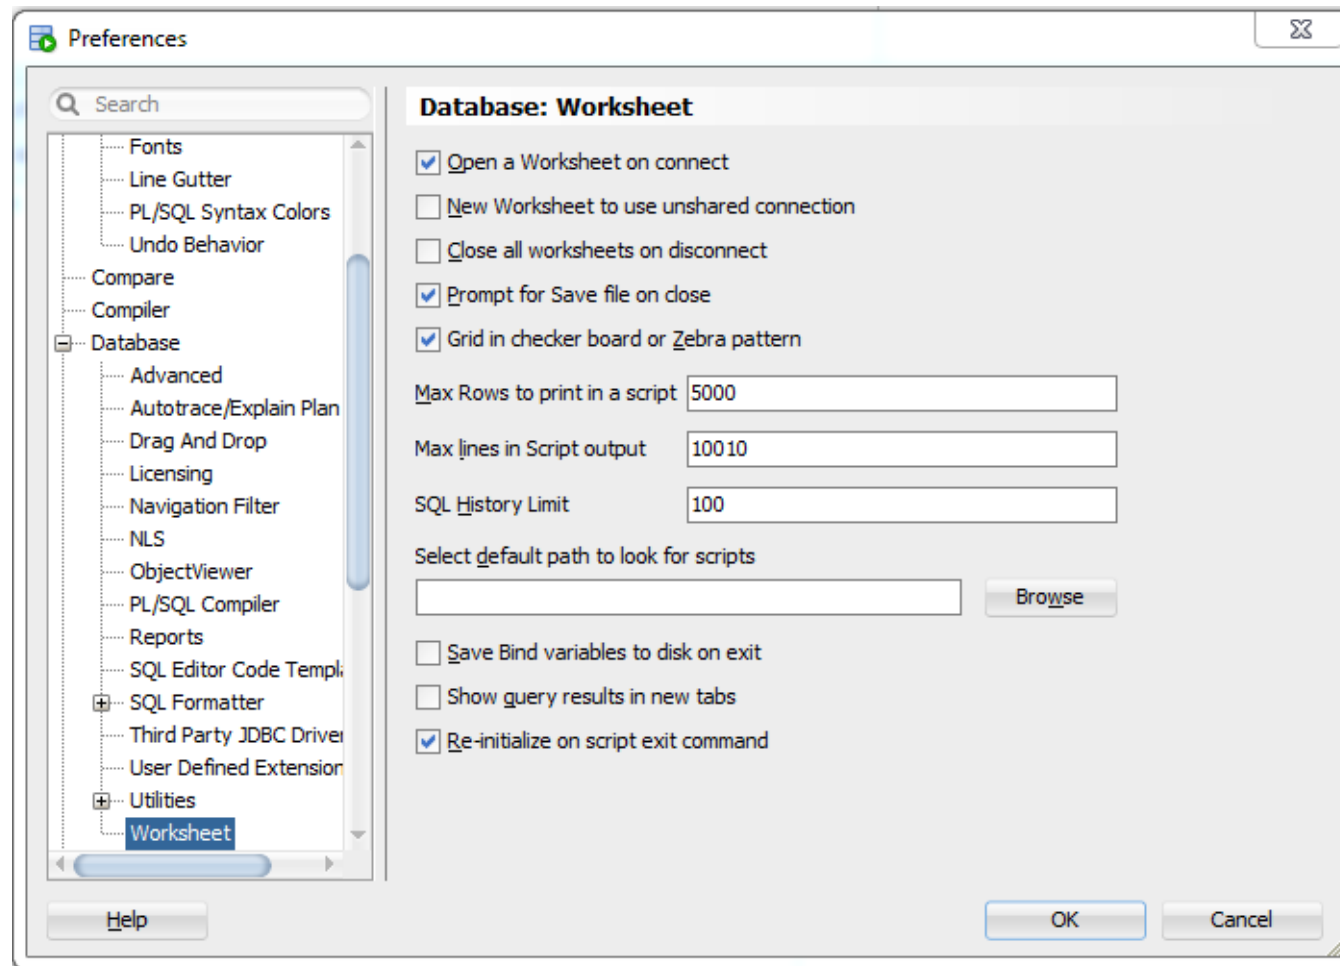
# Line Numbers



# Zebra Pattern Print Output

- You may have noticed that the output in my queries used alternate colours for each line
- It alternated between a light grey and a cream colour
- If this is not done the output appears in one colour, white, still in a grid
- This method just makes it easier to read the output
- Again in the Preferences menu, select Database, scroll down till you see Worksheet, click the option Grid in checker board or Zebra pattern

# Zebra Pattern Print Output



# Summary

- Identify keywords, mandatory clauses, and optional clauses in a SELECT statement
- Select and view all, one or multiple columns of a table
- Use a column alias to clarify contents of a column
- Perform basic arithmetic operations in a SELECT clause
- Remove duplicate items using either DISTINCT or UNIQUE keywords
- Some basic settings for SQL Developer